

# Deep Learning

## lecture 4

### Energy-Based Model

Yi Wu, IIS

Spring 2024

Mar-22

# Logistics

- Coding Project 2 due in 1 week
  - Use local compute for coding & Colab for testing
  - Cloud for long-term training
  - Any questions can be posted in Dingding channel
- Tips
  - Tricks: overfitting & regularization
  - First overfit!
    - **Learning rate decay, architecture, initialization, normalization and preprocess ...**
  - Then regularize!
  - Be aware of your model size and computation (flops)!

# Overview of Lecture 3

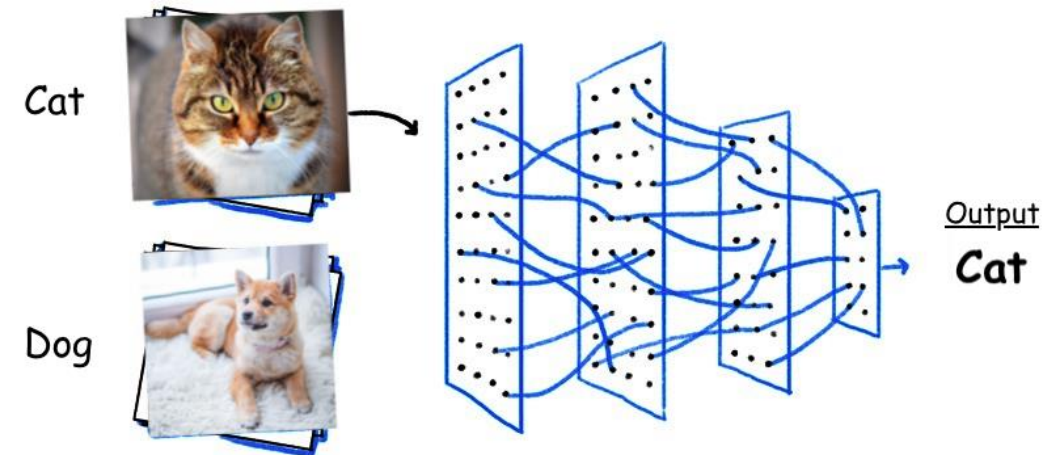
- Algorithm Design
  - Principle 1: adaptive and decoupled learning rate
  - Principle 2: momentum
  - Principle 3: second-order information is great! Let's approximate it
- Practical Algorithms
  - Mini-Batch SGD
  - Momentum SGD
  - AdaGrad, RMSProp, AdaDelta
  - Adam

# Overview of Lecture 3

- Regularizations
  - Goal: stabilize gradients and generalization!
  - Gradient Tricks:
    - Initialization, Gradient Clipping
  - Generalization Tricks:
    - Weight Decay, Dropout, Data Augmentation, Early Stopping
  - Normalization Layers:
    - BatchNorm, LayerNorm
  - Other:
    - Ensemble & practice makes perfect 😊
- Architecture
  - Residual Connection, Dense Connection
  - Fully Convolutional Network


# Story So Far

- History
  - Lecture 1
    - first neural network (1943) to recent advances in deep learning
- Supervised Learning (Classification)
  - Lecture 2
    - MLP and basic components; Backpropagation
  - Lecture 3
    - Algorithms, Tricks and Architecture
- Discriminative Model
  - $P(y|X)$
  - Labeled data;  $X \rightarrow y$



# Afterwards

- What if we want to generate  $X$ ?
- Generative Model
  - $P(X, y) = P(y) * P(X|y)$
  - Or just  $P(X)$
- Lecture 3~7
  - Deep Generative Models
  - Ask the neural network to generate a cat!



# Today's Lecture: Energy-Based Models

- A particularly flexible and general form of ***generative model***
- Part 1: Hopfield Network
  - The simplest model that can memorize and generate patterns
- Part 2: Boltzmann Machine
  - The first deep generative model
- Part 3: General Energy-Based Models


# Today's Lecture: Energy-Based Models

- A particularly flexible and general form of ***generative model***
- Part 1: Hopfield Network
  - The simplest model that can memorize and generate patterns
- Part 2: Boltzmann Machine
  - The first deep generative model
- Part 3: General Energy-Based Models




# Classification

- Recap: Classification
  - Layer-by-layer computation
  - Computation Graph: Directed Acyclic Graph
  - Feedforward networks




- What about ...
  - Loops!



# A Loopy Network

- A “fully-connected” network
  - Each neuron receives inputs from all the other neurons
  - $y_i = +1$  or  $-1$  with hard thresholding

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases} \quad y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$




The output of a neuron affects the input to the neuron

# Hopfield Network

- A “fully-connected” network
  - Each neuron receives inputs from all the other neurons
  - $y_i = +1$  or  $-1$  with hard thresholding
  - Symmetric weights

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases} \quad y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$




A symmetric network:  
 $w_{ij} = w_{ji}$

# Hopfield Network

- A Hopfield Network may not be stable!
  - At each time each neuron receives a “field”  $\sum_{j \neq i} w_{ji} y_j + b_i$
  - If the sign of neuron matches the sign of the field, it flips

$$y_i \leftarrow -y_i \text{ if } y_i \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) < 0$$

- This can further cause other neurons to flip!



$$y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$


# Hopfield Network

- Neurons flip if its sign does not match its local “field”

- $y_i \leftarrow -y_i$  if  $y_i(\sum_{j \neq i} w_{ji}y_j + b_i) < 0$  for all neurons
- Repeat until no neuron flips
- Will this process converge?


$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

$$y_i = \Theta\left(\sum_{j \neq i} w_{ji}y_j + b_i\right)$$



# Hopfield Network

- Let  $y_i^-$  denote the value of  $y_i$  before a “flip”
- Let  $y_i^+$  denote the value of  $y_i$  after a “flip”
- If  $y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) \geq 0$ , nothing happen  
$$y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) - y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) = 0$$




$$y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

# Hopfield Network

- Let  $y_i^-$  denote the value of  $y_i$  before a “flip”
- Let  $y_i^+$  denote the value of  $y_i$  after a “flip”
- If  $y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) \geq 0$ , nothing happen
- If  $y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) < 0$ ,  $y_i^+ = -y_i^-$   
$$y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) - y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) = 2y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$



$$y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$


$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

# Hopfield Network

- Let  $y_i^-$  denote the value of  $y_i$  before a “flip”
- Let  $y_i^+$  denote the value of  $y_i$  after a “flip”
- If  $y_i^- (\sum_{j \neq i} w_{ji} y_j + b_i) \geq 0$ , nothing happen
- If  $y_i^- (\sum_{j \neq i} w_{ji} y_j + b_i) < 0$ ,  $y_i^+ = -y_i^-$

*Every flip increases*  
 $2y_i (\sum_{j \neq i} w_{ji} y_j + b_i)$

$$y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) - y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) = 2y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) \quad \text{Positive!}$$



$$y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$



# Hopfield Network


- Consider the sum over every pair of neurons (assume  $w_{ii} = 0$ )

$$D(y_1, \dots, y_N) = \sum_{i < j} y_i w_{ij} y_j + y_i b_i$$

- Any flip that changes  $y_i^-$  to  $y_i^+$  increases  $D(y_1, \dots, y_N)$

$$\Delta D = D(\dots, y_i^+, \dots) - D(\dots, y_i^-, \dots) = 2y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) > 0$$

- Convergence?



# Hopfield Network


- $D$  is upper-bounded (we only change  $y_i$ )

$$D(y_1, \dots, y_N) = \sum_{i < j} w_{ij} y_i y_j + \sum_i b_i y_i \leq \sum_{i < j} |w_{ij}| + \sum_i |b_i|$$

- $\Delta D$  is lower-bounded

$$\Delta D_{\min} = \min_{i, \{y_j\}} 2 \left| \sum_j w_{ij} y_j + b_i \right| > 0$$

- $\{y_i\}$  converges with a finite number of iterations!
  - $\{y_i\}$ : *state*



# Hopfield Network

- The **Energy** of Hopfield Network


$$E = -D = -\sum_{i < j} w_{ij} y_i y_j - \sum_i b_i y_i$$

- The evolution of Hopfield network always decreases its energy!
- The concept of Energy

- Magnetic dipoles in a disordered magnetic material
- Each dipole tries to align itself to the local field
- Field at a particular dipole  $f(p_i)$ ,  $p_i$  is the position of  $x_i$

$$f(p_i) = \sum_{j \neq i} J_j x_j + b_i$$

- **Ising model** of magnetic materials (Ising and Lenz, 1924)



# Hopfield Network

- Ising model for magnetic materials

- Total field for a dipole

$$f(p_i) = \sum_{j \neq i} J_{ij} x_j + b_i$$

- Response of a dipole

- $x_i \leftarrow -x_i$  if  $\text{sign}(x_i f(p_i)) \neq 1$

- Hamiltonian (total energy) of the system

$$E = -\frac{1}{2} \sum_i x_i f(p_i) = -\sum_{i < j} J_{ij} x_i x_j - \sum_i b_i x_i$$

- The system *evolves* to minimize the energy




- A dipole stop flipping if that flip increases the energy  $\rightarrow$  a local minimum

# Hopfield Network

- The Hopfield network (simplified)

$$E = - \sum_{i < j} w_{ij} y_i y_j$$

- Network evolution arrives at a local optimum in the energy contour
  - Every change in the network state decreases the energy
- Any small jitter from this stable state **returns** it to the stable state






# Hopfield Network

- The Hopfield network (simplified)


$$E = - \sum_{i < j} w_{ij} y_i y_j$$

- Each local optimum state is a “stored” pattern
  - If the network is initialized close to a stored pattern, it evolves to the pattern
- *Associated Memory (content addressable memory)*



# Hopfield Network

- Image Reconstruction by Hopfield Network (1982)



Hopfield network reconstructing degraded images  
from noisy (top) or partial (bottom) cues.


- *How can we store the desired patterns?*

# Hopfield Network: Training


- Let's teach the network to store this image
  - $N$  pixels  $\rightarrow N$  neurons
  - Symmetric weights  $\rightarrow \frac{1}{2}N(N - 1)$  parameters to learn



- Design  $\{w_{ij}\}$  such that the energy is at a local minimum for a desired pattern  $y$ 
  - Redundancy!  $y$  &  $-y$  will be both stored



$$E = - \sum_i \sum_{j < i} w_{ji} y_j y_i$$





# Hopfield Network: Training

- Let's teach the network to store this image
  - $N$  pixels  $\rightarrow N$  neurons
  - Symmetric weights  $\rightarrow \frac{1}{2}N(N - 1)$  parameters to learn
- Design  $\{w_{ij}\}$  such that the energy is at a local minimum for a desired pattern  $y$ 
  - Hebbian Learning Rule  $w_{ij} \leftarrow y_i y_j$  (1949)
  - $E = -\sum_{i < j} w_{ij} y_i y_j = -\frac{1}{2}N(N - 1) \rightarrow$  lowest possible energy!



# Hopfield Network: Training

- What if we want to store **multiple** patterns?
  - $P = \{y^p\}$   $N_p$  patterns
  - Hebbian Learning Rule


$$w_{ij} = \frac{1}{N_p} \sum_p y_i^p y_j^p$$

- The issue of Hebbian Learning
  - Spurious local optima


# Hopfield Network: Training

- Example: 4-dimensional Hopfield Network with Hebbian Learning
  - Two orthogonal patterns to store
    - *Let's assume the value of each neuron is 1 or -1*

***Left:  
desired  
patterns***



**2 nonorthogonal patterns**




***Right:  
stored  
patterns***



# Hopfield Network: Training


- Example: 4-dimensional Hopfield Network with Hebbian Learning
  - Three patterns
  - *How many patterns can a Hopfield network store?*

*Left:  
desired  
patterns*



**Disappeared!**

3 nonorthogonal patterns




**"Fake" memory**

*Right:  
stored  
patterns*

# Hopfield Network: Training

- How many patterns can a Hopfield network store?
- A fact: you can store all the  $2^N$  patterns!
- Solution: find any  $N$  orthogonal patterns
  - Prove this fact in your homework 😊

*A “flatten” landscape does not help evolve desired patterns!*



# Hopfield Network: Training

- We want to construct a network with desired **stable local optimum**
  - A pattern can be recovered after 1-bit change
- For a specific set of  $K$  patterns, we can always build a network for which all patterns are stable provided  $K \leq N$ 
  - Mostafa and St. Jacques (1985)
  - For large  $N$ , the upper bound on  $K$  is actually  $\frac{N}{4} \log N$ 
    - McElice et. al. (1987)
  - Still possible with undesired local minimum
- **How can we find the weights?**
  - $K$  patterns remembered
  - Avoid undesired local minimum as much as we can

# Hopfield Network: Optimization

- Problem Formulation

- Desired patterns  $P = \{y^p\}$
- Energy function  $E(y) = -\frac{1}{2}y^T W y$  (we omit bias for simplicity)

- Objective for  $W$

- Minimize  $E$  for all  $y^p$
- It should also maximize  $E$  for all non-desired patterns!

$$W = \arg \min_W \sum_{y \in P} E(y) - \sum_{y' \notin P} E(y')$$


- Gradient Descent

$$W \leftarrow W - \eta \left( \sum_{y \in P} y y^T - \sum_{y' \notin P} y' y'^T \right)$$

# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P} y'y'^T \right)$$






# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P} y'y'^T \right)$$

- The first term is minimizing the energy of desired patterns!




# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P} y'y'^T \right)$$

- The second term essentially raises all the patterns in the space

- **Issue??**




# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ \& } y' \in \textit{Valley}} y'y'^T \right)$$

- Let's just focus on the valleys!




# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ \& } y' \in \textit{Valley}} y'y'^T \right)$$

- Let's just focus on the valleys!
- **But how can we find the valleys?**




# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ \& } y' \in \textit{Valley}} y'y'^T \right)$$

- Let's just focus on the valleys!
- But how can we find the valleys?
- **Evolution of Hopfield Network will converge to a valley**



# Hopfield Network: Optimization

- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ \& } y' \in \textit{Valley}} y'y'^T \right)$$

- Compute outer-products of desired patterns  $y$
- Randomly initialize  $y'$  for multiple times
  - Run evolution for random  $y'$  until convergence
  - Calculate outer-product of  $y'$
- Compute gradient and update  $W$

# Hopfield Network: Optimization


- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ \& } y' \in \textit{Valley}} y'y'^T \right)$$

- Compute outer-products of desired patterns  $y$
- **Randomly initialize**  $y'$  for multiple times
  - Run evolution for random  $y'$  until convergence
  - Calculate outer-product of  $y'$
- Compute gradient and update  $W$
- **Valleys are NOT equivalently important...**

# Hopfield Network: Optimization


- Which valleys are important?
- Primary object: ensure desired patterns stable
  - We want to ensure desired patterns are in broad valleys






# Hopfield Network: Optimization

- Which valleys are important?
- Primary object: ensure desired patterns stable
  - We want to ensure desired patterns are in broad valleys
  - **Spurious valleys around desired patterns are more important to eliminate**



# Hopfield Network: Optimization

- Which valleys are important?
- Primary object: ensure desired patterns stable
  - We want to ensure desired patterns are in broad valleys
  - Spurious valleys around desired patterns are more important to eliminate
  - **Evolution from desired patterns**



# Hopfield Network: Optimization


- Update rule for  $W$

$$W \leftarrow W - \eta \left( \sum_{y \in P} yy^T - \sum_{y' \notin P \text{ \& } y' \in Valley} y'y'^T \right)$$

- Compute outer-products of desired patterns  $y$
- Initialize  $y'$  by all the desired patterns
  - Run evolution for random  $y'$  until convergence
  - Calculate outer-product of  $y'$
- Compute gradient and update  $W$
- Still issues?


# Hopfield Network: Optimization

- Recap: we raise the valleys next to the desired patterns




# Hopfield Network: Optimization

- Recap: we raise the valleys next to the desired patterns
- What if a pattern is close to the valley?
  - Naively forcing a valley to raise may hurt the learned representation
  - Particularly challenging when  $y$  are continuously valued (e.g., tanh activation)




# Hopfield Network: Optimization

- New idea: we only raise the “neighborhood” of desired patterns!
  - It is sufficient to make each desired pattern a valley
  - Note: we want to raise the “decent” neighborhood



# Hopfield Network: Optimization

- New idea: we only raise the “neighborhood” of desired patterns!
  - It is sufficient to make each desired pattern a valley
  - Note: we want to raise the “decent” neighborhood
- Implementation
  - We initialize  $y'$  by the desired patterns
  - **Only perform evolution for a few steps!**



# Hopfield Network: SGD Optimization

- SGD update rule for  $W$


$$W \leftarrow W - \eta (E_{y \in P} [yy^T] - E_{y'} [y'y'^T])$$

- Compute outer-products of random desired pattern  $y$
- Initialize  $y'$  by a random desired pattern
  - Run evolution for random  $y'$  for a few timesteps (2~4)
  - Calculate outer-product of  $y'$
- Compute gradient and update  $W$
- In theory,  $O(N)$  patterns can be stored in the network (with undesired valleys)
  - How to store more patterns?



# The Expanded Network

- Idea: introduce redundant neurons to increase network capacity
- Original  $N$  neurons for patterns: visible neurons
- Additional  $K$  neurons: hidden neurons




# The Expanded Network

- Idea: introduce redundant neurons to increase network capacity
- Original  $N$  neurons for patterns: visible neurons
- Additional  $K$  neurons: hidden neurons



# The Expanded Network


- $N$  dimensional pattern  $\rightarrow N + K$  dimension
  - How can we store the patterns with  $K$  additional units?
  - Possible solution: random filling of  $K$  units (not great but okay...)



# The Expanded Network


- How to retrieve the stored pattern?
  - Still evolution?
  - Evolution is performed on the entire network but we only care about  $N$  units

A mechanism that can decouple  $N$  visible units when needed!




# The Expanded Network

- How to retrieve the stored pattern?
  - Idea: Probabilistic Framework  $P_w(v, h)$  **Let's borrow some ideas from physics!**
    - Consider desired patterns by computing the marginal  $P_w(v) = \sum_h P(v, h)$
    - How to convert Hopfield network to a distribution?



# The Helmholtz Free Energy of a System

- Recap: A thermodynamic (热力学) system
  - We previously discussed a discrete-time version
  - In fact, it is a probabilistic system
- A thermodynamic system at temperature  $T$ 
  - $P_T(S)$  the probability of the system at state  $S$
  - $E_T(S)$  the potential energy at state  $S$
  - $U_T$  the internal energy, the capability to do work
  - $H_T$  the entropy, internal disorder of the system
  - $k$  Boltzmann constant
  - Free energy  $F_T = U_T - kTH_T$





# The Helmholtz Free Energy of a System

- A thermodynamic system at temperature  $T$ 
  - Internal energy  $U_T = \sum_S P_T(S) E_T(S)$
  - Entropy  $H_T = - \sum_S P_T(S) \log P_T(S)$
  - Free energy  $F_T = \sum_S P_T(S) E_T(S) - kT \sum_S P_T(S) \log P_T(S)$
- Minimum Free-Energy Principle
  - A system held at temperature  $T$  anneals by varying the rate at which it visits the various states until a minimum free-energy state is achieved
- Boltzmann Distribution
  - The probability distribution of states at equilibrium

# The Helmholtz Free Energy of a System

- Free energy

$$F_T = \sum_S P_T(S) E_T(S) + kT \sum_S P_T(S) \log P_T(S)$$

- Boltzmann distribution (minimize  $F_T$  w.r.t.  $P_T(S)$ )

$$P_T(S) = \frac{1}{Z} \exp\left(-\frac{E_T(S)}{kT}\right)$$

- It is also known as Gibbs distribution
- $Z$  normalizing constant

Given an energy function  $E_T(S)$ , if we follow a proper physical evolution process, the system will converge to the Boltzmann distribution



# Stochastic Hopfield Network

- Let's model our Hopfield network as a thermodynamic system

- $T = k = 1$  for simplicity
- Energy


$$E(y) = - \sum_{i < j} w_{ij} y_i y_j - b_i y_i$$

- Boltzmann Probability


$$P(y) = \frac{1}{Z} \exp \left( \sum_{i < j} w_{ij} y_i y_j + b_i y_i \right)$$

- Stochastic Hopfield Network

- Models the stationary probability distribution of states
- Generative model: generate state from  $P(y)$




# Stochastic Hopfield Network



- Let's consider the “flip” operation
  - Deterministic  $\rightarrow$  probabilistic
  - Goal: change  $y_i$  to 1 with probability  $P(y_i = 1|y_{j \neq i})$
- Assume  $y$  and  $y'$  only differ at position  $i$  and  $y'_i = -1$ 
  - $\log P(y) = -E(y) + C$
  - $E(y) = -\sum_{i < j} w_{ij} y_i y_j - b_i y_i$
  - $\log P(y) - \log P(y') = E(y') - E(y) = -\sum_j w_{ij} y_j - 2b_i$
  - $\log \frac{P(y)}{P(y')} = \log \frac{P(y_i = 1|y_{j \neq i})P(y_{j \neq i})}{P(y'_i = -1|y'_{j \neq i})P(y'_{j \neq i})} = \log \frac{P(y_i = 1|y_{j \neq i})}{1 - P(y_i = 1|y_{j \neq i})} = -\sum_j w_{ij} y_j - 2b_i$


# Stochastic Hopfield Network



- Let's consider the “flip” operation
  - Deterministic  $\rightarrow$  probabilistic
  - Goal: change  $y_i$  to 1 with probability  $P(y_i = 1|y_{j \neq i})$
- Assume  $y$  and  $y'$  only differ at position  $i$  and  $y'_i = -1$ 
  - $\log P(y) = -E(y) + C$
  - $E(y) = -\sum_{i < j} w_{ij} y_i y_j - b_i y_i$
  - $\log P(y) - \log P(y') = E(y') - E(y) = -\sum_j w_{ij} y_j - 2b_i$
  - $\log \frac{P(y)}{P(y')} = \log \frac{P(y_i = 1|y_{j \neq i})P(y_{j \neq i})}{P(y'_i = -1|y'_{j \neq i})P(y'_{j \neq i})} = \log \frac{P(y_i = 1|y_{j \neq i})}{1 - P(y_i = 1|y_{j \neq i})} = -\sum_j w_{ij} y_j - 2b_i$
- A sigmoid conditional:  $P(y_i = 1|y_{j \neq i}) = \frac{1}{1 + \exp(-\sum_j w_{ij} y_j - 2b_i)}$

*This is also called Gibbs sampling (remember the name for now 😊)*

# Stochastic Hopfield Network



- The whole update rule
  - Field at  $y_i$ :  $z_i = \sum_j w_{ij}y_j + 2b_i$
  - $P(y_i = 1|y_{j \neq i}) = \frac{1}{1+\exp(-z_i)} = \sigma(z_i)$
- Running the network
  - Randomly initialize  $y$
  - Cycle over  $y_i$ , fixed other variables fixed and sample  $y_i$  according to the conditional probability
  - After “convergence”, we can get samples of  $y$  according to  $P(y)$
  - *This sampling procedure is called Gibbs sampling*
  - **How can we retrieve the stored pattern???**
    - **This is a stochastic process!**


Field quantifies the  
delta energy of flip

# Stochastic Hopfield Network

- Network evolution
  - initialize  $y_0$
  - For  $1 \leq i \leq N$ ,  $y_i(t + 1) \sim \text{Bernoulli}(\sigma(z_i(t)))$
  - Until convergence
- Retrieve a stored pattern  $y$ 
  - Given sequence of samples  $y_0, \dots, y_L$
  - Simply take the average of final  $M$  samples


$$y_i = I \left[ \frac{1}{M} \sum_{t=L-M+1}^L y_i(t) > 0 \right]$$

- If you want a probability instead of a single vector, you can use the frequency derived from  $\{y_{L-M+1}, \dots, y_L\}$  to approximate the stationary distribution
- **In many applications, we simply take  $M = 1$  (output  $y_L$ )**



# Stochastic Hopfield Network: Annealing


- Find the state with lowest energy?
- Network evolution
  - initialize  $y_0, T \leftarrow T_{\max}$
  - Repeat
    - Repeat a few cycles
      - For  $1 \leq i \leq N, y_i(T) \sim \text{Bernoulli} \left( \sigma \left( \frac{1}{T} z_i(T) \right) \right)$
    - $y_i(\alpha T) \leftarrow y_i(T); T \leftarrow \alpha T$
  - Until convergence
- Final state as the retrieved pattern
  - With temperature annealing, the system will converge to the most likely state
  - Possibly local minimum in practice



# Boltzmann Machine

- A generative Model
  - $E(y) = -\frac{1}{2}y^T W y$
  - $P(y) = \frac{1}{Z} \exp\left(-\frac{E(y)}{T}\right)$
  - Parameter  $W$
- It has a probability for producing any binary pattern  $y$ 
  - We assume  $y_i = 0$  or  $1$  (or  $\pm 1$ )

**How to learn  $W$  for desired patterns?**



$$z_i = \frac{1}{T} \sum_j w_{ji} s_j$$

$$P(s_i = 1 | s_{j \neq i}) = \frac{1}{1 + e^{-z_i}}$$

# Boltzmann Machine: Training

- Goal
  - Remember a set of desired patterns  $P = \{y^p\}$
  - Now we have a probability distribution
- Objective: maximum likelihood learning (assume  $T = 1$ )
  - Probability of a particular pattern

$$P(y) = \frac{\exp\left(\frac{1}{2} y^T W y\right)}{\sum_{y'} \exp\left(\frac{1}{2} y'^T W y'\right)}$$

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp\left(\frac{1}{2} y'^T W y'\right)$$



# Boltzmann Machine: Training

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp \left( \frac{1}{2} y'^T W y' \right)$$

- Gradient Ascent  $\nabla_{w_{ij}} L$

# Boltzmann Machine: Training

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp \left( \frac{1}{2} y'^T W y' \right)$$

- Gradient Ascent  $\nabla_{w_{ij}} L$

- $\nabla_{w_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j$

# Boltzmann Machine: Training

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp \left( \frac{1}{2} y'^T W y' \right)$$

- Gradient Ascent  $\nabla_{w_{ij}} L$

$$\nabla_{w_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \sum_{y'} \frac{\exp \left( \frac{1}{2} y'^T W y' \right)}{Z} \cdot y'_i y'_j \quad \text{Exponentially many terms!}$$

# Boltzmann Machine: Training

- Maximize log-likelihood

$$L(W) = \frac{1}{N_P} \sum_{y \in P} \frac{1}{2} y^T W y - \log \sum_{y'} \exp \left( \frac{1}{2} y'^T W y' \right)$$

- Gradient Ascent  $\nabla_{w_{ij}} L$

- $\nabla_{w_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \sum_{y'} \frac{\exp\left(\frac{1}{2} y'^T W y'\right)}{Z} \cdot y'_i y'_j$
- $\nabla_{w_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \mathbb{E}_{y'}[y'_i y'_j]$  **Monte-Carlo Approximation**
- Draw a set of samples  $S$  for  $y'$  according to the probability,
- $\nabla_{w_{ij}} L = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \frac{1}{|S|} \sum_{y' \in S} y'_i y'_j$

# Boltzmann Machine: Training

- Maximize log-likelihood with  $M$  Monte-Carlo samples

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \frac{1}{M} \sum_{y' \in S} y'_i y'_j$$

- How to draw samples from  $P(y)$ ?
  - Running the stochastic network (Gibbs sampling)
  - Randomly initialize  $y(0)$
  - Cycle over  $y_i(t)$ , sampling according to  $P(y_i(t) | y_{j \neq i}(t))$
  - After convergence, we get a sequence of samples  $\{y(0), \dots, y(L)\}$
  - Get the final  $M$  states as samples  $S = \{y(L - M + 1), \dots, y(L)\}$

# Boltzmann Machine: Training

- Overall Training


- Initialize  $W$
- Maximize log-likelihood with  $M$  Monte-Carlo samples

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P} \sum_{y \in P} y_i y_j - \frac{1}{M} \sum_{y' \in S} y'_i y'_j$$

- $w_{ij} \leftarrow w_{ij} + \eta \nabla_{w_{ij}} L(W)$  (we are maximizing likelihood)


- Remark

- We can also run Gibbs sampling from states in  $P$  (will discuss later...)



# Boltzmann Machine with Hidden Neurons

- Let's get back to hidden neurons!
  - $v$  visible neurons (pattern),  $h$  hidden neurons
  - $y = (v, h)$
- A joint probability distribution
  - $P(y) = P(v, h)$
  - $P(v) = \sum_h P(v, h)$ 
    - **The marginal distribution!**
  - $h$ : latent representation
- New objective
  - Maximize the marginal probability



# Boltzmann Machine with Hidden Neurons

- Maximum log-likelihood learning

$$P(v) = \sum_h P(v, h) = \sum_h \frac{\exp(y^T W y)}{\sum_{y'} \exp(y'^T W y')}$$

$$L(W) = \frac{1}{|P|} \sum_{v \in P} \log \left( \sum_h \exp(y^T W y) \right) - \log \left( \sum_{y'} \exp(y'^T W y') \right)$$

- Gradient  $\nabla L(W)$ ?



# Boltzmann Machine with Hidden Neurons

- Maximum log-likelihood learning

$$P(v) = \sum_h P(v, h) = \sum_h \frac{\exp(y^T W y)}{\sum_{y'} \exp(y'^T W y')}$$

$$L(W) = \frac{1}{|P|} \sum_{v \in P} \log \left( \sum_h \exp(y^T W y) \right) - \log \left( \sum_{y'} \exp(y'^T W y') \right)$$

- Gradient  $\nabla L(W)$ ?

**Monte-Carlo Estimate!**

# Boltzmann Machine with Hidden Neurons

- Maximum log-likelihood learning

$$P(v) = \sum_h P(v, h) = \sum_h \frac{\exp(y^T W y)}{\sum_{y'} \exp(y'^T W y')}$$

$$L(W) = \frac{1}{|P|} \sum_{v \in P} \log \left( \sum_h \exp(y^T W y) \right) - \log \left( \sum_{y'} \exp(y'^T W y') \right)$$

- Gradient  $\nabla L(W)$ ?
  - The first term is also in the form of log-sum
  - Monte Carlo Estimate for each  $v \in P$ !

# Boltzmann Machine with Hidden Neurons

- Maximum log-likelihood learning

$$\nabla_{w_{ij}} L(W) = \frac{1}{|P|} \sum_{v \in P} E_h[y_i y_j] - E_{y'}[y'_i y'_j]$$

- Second term
  - Freely generate samples w.r.t.  $p(y)$
  - Random initialization, cyclic Gibbs sampling
- First term
  - Generate samples w.r.t.  $p(y)$  conditioned on a fixed  $v$
  - Randomly initialize  $h$ , run Gibbs sampling over  $h$


# Boltzmann Machine with Hidden Neurons

- Overall Training

- Initialize  $W$
- For  $v \in P$ , fixed the visible neurons, run Gibbs sampling to get  $K$  samples
  - Collect all conditioned samples as  $S_c$
- Randomly initialize all neurons, run Gibbs sampling to get  $M$  samples
  - Collect free samples as  $S$
- Maximize log-likelihood with  $M$  Monte-Carlo samples

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P K} \sum_{y \in S_c} y_i y_j - \frac{1}{M} \sum_{y' \in S} y'_i y'_j$$

- $w_{ij} \leftarrow w_{ij} + \eta \nabla_{w_{ij}} L(W)$




# Boltzmann Machine

- Summary

- A stochastic version of Hopfield Network
- Nice mathematical properties
- Large capacity for storing patterns (with hidden neurons)
- Pattern generation
  - Gibbs sampling
- Pattern completion
  - Conditioned Gibbs sampling

- **Classification??**

- $y = (v, h, c)$ ,  $c$  is label
- $c$  as a one-hot vector (0-1 variables)
- Posterior  $P(c|v)$
- Even conditional generation:  $P(v|c)$ !




Hopfield network reconstructing degraded images  
from noisy (top) or partial (bottom) cues.

# Boltzmann Machine

- The issue
  - Training is hard!
  - Gibbs sampling may take a very long time to converge
    - also called ***mixing-time***
  - Not really applicable for large problems
- Can we design a better structure for faster Gibbs sampling mixing?

# Restricted Boltzmann Machine

- A particularly structured Boltzmann Machine
  - A partitioned structure
  - Hidden neurons are only connected to visible neurons
  - No intra-layer connections
  - *Invented under the name Harmonium by Paul Smolensky in 1986*
  - *Became promise after Hinton invented fast learning algorithms in mid-2000*



# Restricted Boltzmann Machine

- Computation Rules: same as Boltzmann machine

- Hidden neurons  $h_i$


$$z_i = \sum_j w_{ij} v_j, \quad P(h_i = 1 | v_j) = \frac{1}{1 + \exp(-z_i)}$$

- Visible neurons  $v_j$

$$z_j = \sum_i w_{ij} h_i, \quad P(v_j = 1 | h_i) = \frac{1}{1 + \exp(-z_j)}$$



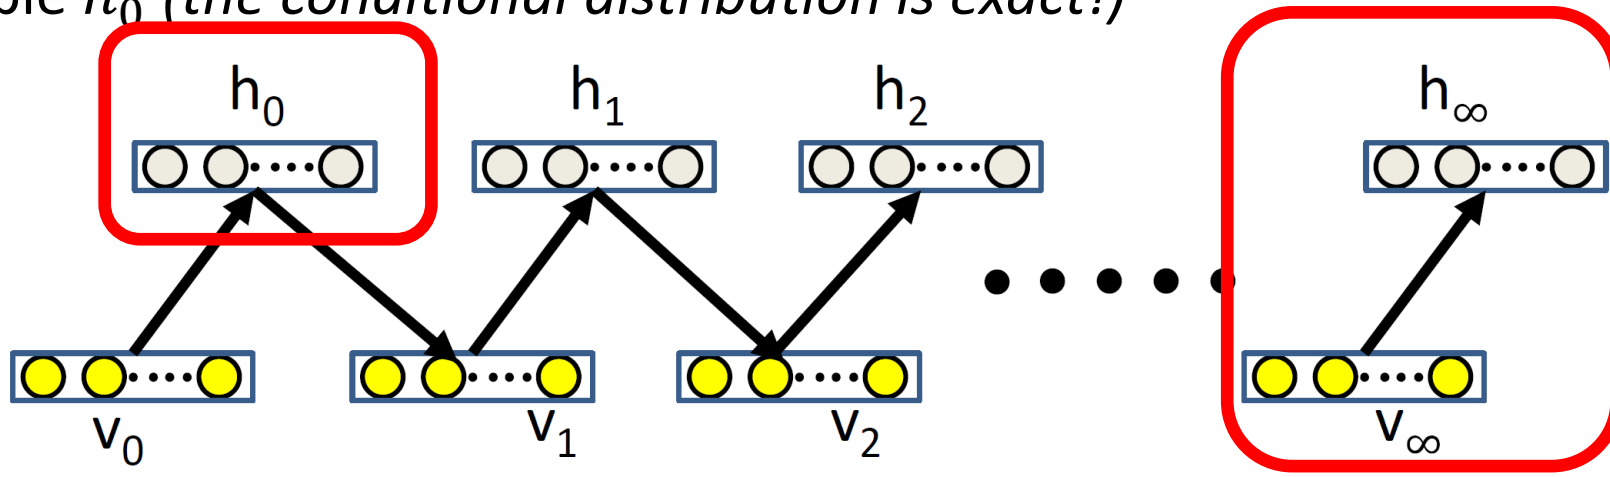
Iterative Sampling!





# Restricted Boltzmann Machine

- Sampling
  - Randomly initialize visible neurons  $v_0$
  - Iterative between hidden neurons and visible neurons
  - Get final sample  $(v_\infty, h_\infty)$
- Conditioned sampling?
  - Initialize  $v_0$  as the desired pattern
  - Sample  $h_0$  (*the conditional distribution is exact!*)




# Restricted Boltzmann Machine

- Maximum Likelihood Estimate

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P K} \sum_{v \in P} v_{0i} h_{0j} - \frac{1}{M} \sum v_{\infty i} h_{\infty j}$$

- No need to lift up the entire energy landscape! (recap)




# Restricted Boltzmann Machine

- Maximum Likelihood Estimate

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P K} \sum_{v \in P} v_{0i} h_{0j} - \frac{1}{M} \sum v_{\infty i} h_{\infty j}$$

- No need to lift up the entire energy landscape! (recap)
  - Raising the neighborhood of desired patterns will be sufficient




# Restricted Boltzmann Machine

- Maximum Likelihood Estimate

$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P K} \sum_{v \in P} v_{0i} h_{0j} - \frac{1}{M} \sum v_{\infty i} h_{\infty j}$$

- No need to lift up the entire energy landscape!
  - One Gibbs sampling will be sufficient




# Restricted Boltzmann Machine

- Maximum Likelihood Estimate


$$\nabla_{w_{ij}} L(W) = \frac{1}{N_P} \sum_{v \in P} v_{0i} h_{0j} - v_{1i} h_{1j}$$

- Only 3 Gibbs sampling steps are needed!
- We can also extend (R)BM to continuous values!
  - If we can explicitly sample from  $P(y_i | y_{j \neq i})$
  - Exponential family! (FYI 😊)
    - “Exponential Family Harmoniums with an Application to Information Retrieval”, Welling et al., 2004



# Deep Boltzmann Machine



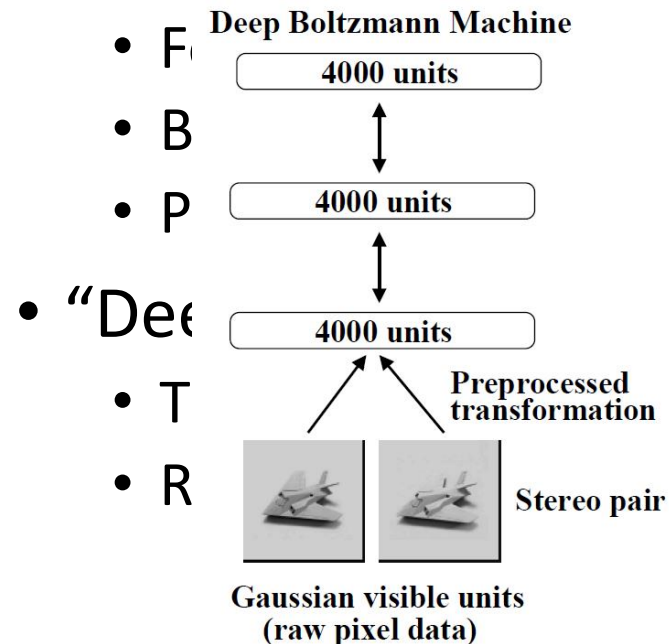
- Can we have a **deep** version of RBM?
  - Deep Belief Net (2006)
  - Deep Boltzmann Machine (2009)
- Sampling?
  - Forward pass: bottom-up
  - Backward pass: top-down
  - Practical Trick: Layer-by-layer pretraining
- “Deep Boltzmann Machine”, AISTATS 2009
  - The very first deep generative model
  - Ruslan Salakhutdinov & Geoffrey Hinton



# Deep Boltzmann Machine

- Can we have a **deep** version of RBM?
  - Deep Belief Net (2006)
  - Deep Boltzmann Machine (2009)

- Sampling?



# Summary

- Hopfield Network
  - The very first generative neural network
- Boltzmann Machine
  - A stochastic version of Hopfield network
  - An undirected probabilistic model
- Restricted Boltzmann Machine
  - Layered structure for fast inference
- Next
  - General formulation of energy-based models



# Energy-Based Model

- Goal of generative model
  - A probability distribution of “patterns”  $P(x)$
- Requirement
  - $P(x) \geq 0$  (non-negative)
  - $\int_x P(x) dx = 1$  (sum to 1)
- Energy-Based Model
  - Energy function:  $E(x; \theta)$  parameterized by  $\theta$
  - $P(x) = \frac{1}{Z} \exp(-E(x; \theta))$
  - $Z = \int_x \exp(-E(x; \theta)) dx$  *partition function*

Why use  $\exp()$  function?  
e.g.  $|x|$  or  $|x|^2$

# Energy-Based Model

- A particular class of density function

$$P(x) = \frac{1}{Z} \exp(-E(x; \theta))$$

- Pros

- Compatible with log-probability measure to capture large variations
- Exponential family (e.g., Gaussian)
- Common in statistical physics
- Extremely flexible, i.e., use any  $E(x)$  you like (e.g., any  $f(x): \mathbb{R}^d \rightarrow \mathbb{R}$ , even CNNs)

- Cons

- Non-trivial to sample and train due to the partition function  $Z$
- Is it possible to avoid  $Z$ ?

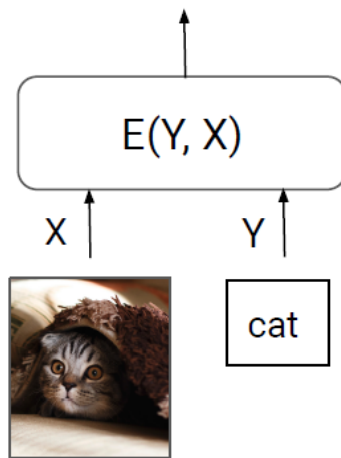
# Energy-Based Model

- A particular class of density function


$$P(x) = \frac{1}{Z} \exp(-E(x; \theta))$$

- The ratio of two samples does not require  $Z$ !


$$\frac{P(x)}{P(x')} = \exp(-E(x; \theta) + E(x'; \theta))$$



*object recognition*



*sequence labeling*



*image restoration*

# Energy-Based Model: Training

- A particular class of density function

$$P(x) = \frac{1}{Z} \exp(-E(x; \theta))$$

- Maximum Likelihood Training

- $L(\theta) = \log P(x) = -E(x; \theta) - \log Z(\theta)$
- use Monte-Carlo Estimates for  $Z(\theta)$

- Contrastive Divergence Algorithm

- $\nabla_{\theta} L(\theta) \approx \nabla_{\theta} (-E(x_{train}; \theta) + E(x_{sample}; \theta))$

- How to sample from an energy-based model?


# A Generic Solution

- Sampling from the energy-based model  $p(s) = \frac{1}{Z} \exp(-E(s))$ 
  - Random initialize  $s^0$
  - $s' \leftarrow s^t + \text{noise}$
  - If  $E(s') < E(s^t)$ ; then accept  $s^{t+1} \leftarrow s'$
  - Else accept  $s'$  with probability  $\exp(E(s^t) - E(s'))$
  - Repeat
- Then after enough iterations, we get samples from  $p(s)$
- *Details to be explained in the next lecture! ☺*

# Modern Energy-Based Model Examples



Du & Mordatch, 2019



Song et. al., 2021

## Additional Readings

OpenAI Blog: <https://openai.com/blog/energy-based-models/>

A nice overview from Yang Song & Diederik Kingma: <https://arxiv.org/abs/2101.03288>

# Summary

- Hopfield Network
  - The first generative neural network
  - Undirected complete graph
- Boltzmann Machine
  - A probabilistic interpretation of Hopfield Network
  - The first deep generative model
- Energy-Based
  - Extremely flexible and powerful, designed to be multi-modal
  - Hard to sample and learn
  - Closely related to probabilistic inference and Bayesian methods

Thanks!