

# Deep Learning

## lecture 6

# Variational Autoencoder

Yi Wu, IIIS

Spring 2024

Apr-7

# Logistics


- Coding Project 2 grading is finished
  - Many submissions have run-time errors during evaluation!
  - Pay attention to format and evaluation!!!
- Coding Project 3 will be released tomorrow
  - Due in 3 weeks
  - Try to start early
- Don't forget about your homework!
  - No late submission

# Today's Topic

- Latent Variable Model
  - Variational inference
- Variational Autoencoder


# Discriminative v.s. Generative (Recap)

- Discriminative model (lecture 2-3)
  - Feedforward networks
  - straightforward to learn
- Generative model
  - The problem itself is hard (need to model high-dimensional data distribution)
  - Inference is non-trivial (posterior distribution)




# Discriminative v.s. Generative (Recap)

- Discriminative model (lecture 2-3)
  - Feedforward networks
  - straightforward to learn
  - Typically require labels (supervised learning)
- Generative model
  - The problem itself is hard (need to model high-dimensional data distribution)
  - Inference is non-trivial (posterior distribution)




# Discriminative v.s. Generative (Recap)

- Discriminative model (lecture 2-3)
  - Feedforward networks
  - straightforward to learn
  - Typically require labels (supervised learning)
- Generative model
  - The problem itself is hard (need to model high-dimensional data distribution)
  - Inference is non-trivial (posterior distribution)
  - **We have a probability distribution to draw samples!**
    - Unsupervised by nature (directly learn  $p(x)$ )
    - Fill missing information (inpainting, learn complex latent structures)






# Generative Model

- Goal: learn  $p(x; \theta)$
- What we have learned ...
  - Energy-based model (lecture 4)
    - $p(x) = \frac{1}{Z} \exp(-E(x))$  (extremely flexible)
    - Sampling: MCMC
      - Gradients! (Stochastic Gradient MCMC)
    - $Z$ : partition function (key challenge)
      - No closed-form density for  $p(x) \rightarrow$  NO MLE Learning!
    - Learning: Contrastive Divergence
      - Decrease  $E(x)$  on data samples & increase  $E(x')$  on non-data samples



# Generative Model

- Goal: learn  $p(x; \theta)$
- What we have learned ...
  - Energy-based model (lecture 4)
    - Most flexible! Hard to sample & learn!
  - Flow model (lecture 5)
    - $x = f(z; \theta)$  where  $f(\cdot; \theta)$  is bijection,  $z \sim N(0, I)$ 
      - $z$  is also called latent representation of  $x$
    - Sampling is straightforward!
    - MLE training is easy
      - $\log p(x) = \log p(z) - \sum_i \log \det |\partial f_i / \partial x|$
      - $f_i$  needs to have a structured Jacobian






# Generative Model

- Goal: learn  $p(x; \theta)$
- What we have learned ...
  - Energy-based model (lecture 4)
    - No explicit sampling  $\rightarrow$  hard training and expensive generation
  - Flow Model (lecture 5)
    - $x = f(z)$ : easy sampling and tractable likelihood
    - Most limited modeling capacity
      - ... **because of the bijection constraint!**
- What if  $x = f(z)$  is NOT a bijection?
  - Still easy generation!
  - What about MLE training? How to compute  $p(x)$ ?

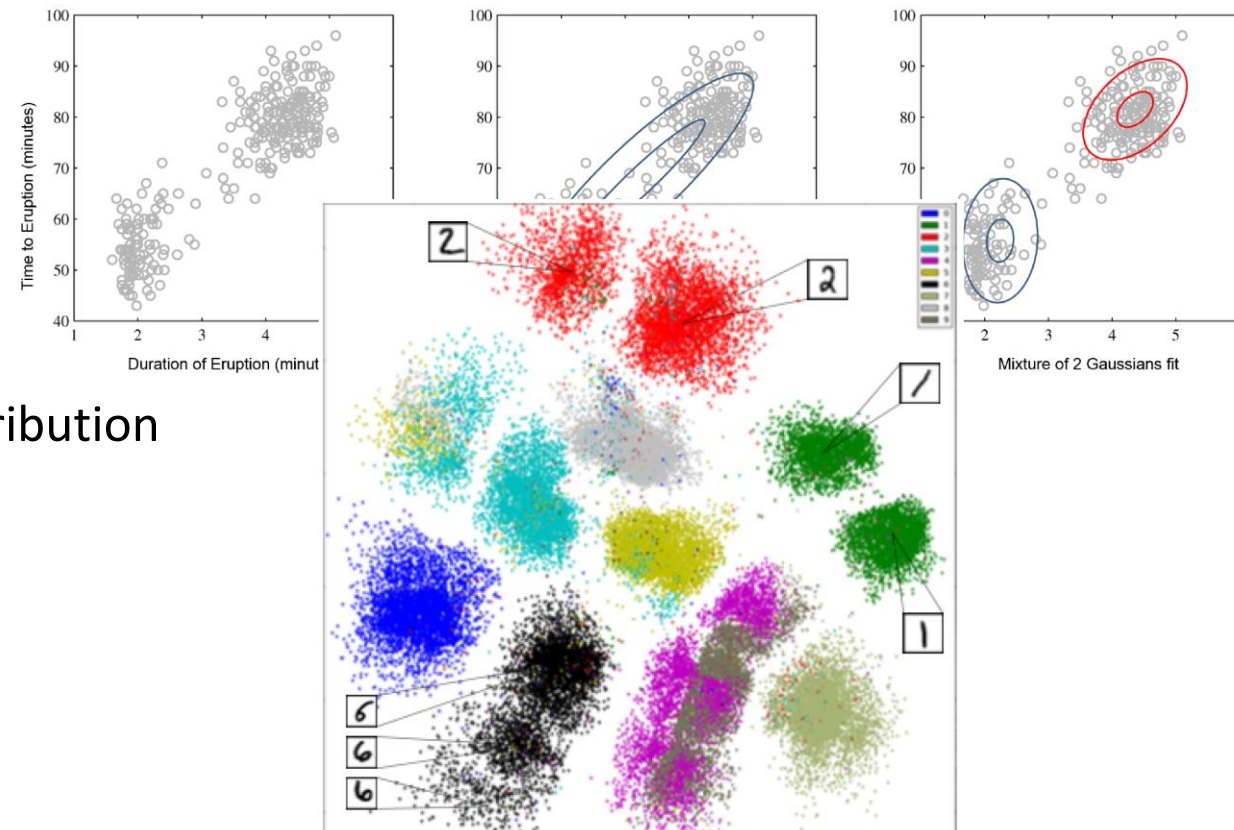

# Latent Variable Model

- A more general formulation:  $p(x, z) = p(z)p(x|z)$ 
  - $x$  data;  $z$  latent variable
  - When  $z$  is given,  $p(x|z)$  is easy to compute
- Example
  - $x$ : image (pixel values)
  - $z$ : latent feature/factors
  - Only grey circle is observed
  - $p(z)$ : prior distribution of factors
  - $p(x|z)$ : a Gaussian/Categorical on each pixel value
    - $p(x|z) = N(\mu(z), \Sigma(z))$



# Latent Variable Model

- $p(x, z) = p(z)p(x|z)$ 
  - $x$  data;  $z$  latent variable
- Example: Gaussian Mixture Model
  - $z \sim \text{Categorical}(w_1, \dots, w_K)$
  - $x \sim N(\mu_z, \Sigma_z)$
  - Generative process
    - Pick a cluster  $z$
    - Generate  $x$  according to the cluster distribution
  - Unsupervised learning
    - Unlabeled data
    - E.g. clustering of handwritten digits



# Latent Variable Model: Training

- Learning the latent variable model
  - Joint probability:  $p(x, z; \theta)$  for random variable  $X$  and  $Z$ 
    - $p(x, z; \theta) = p(z; \theta)p(x|z; \theta)$
  - Dataset  $D = \{x^{(i)}\}$  for  $X$ , variable  $Z$  is never observed
- Maximal Likelihood Learning

$$L(\theta) = \log \prod_{x \in D} p(x; \theta) = \sum_{x \in D} \log \sum_z p(x, z; \theta)$$

- Marginal probability can be expensive to compute!
  - When  $z$  is continuous, the objective even becomes intractable

# Latent Variable Model: Training

- Learning the latent variable model
  - Joint probability:  $p(x, z; \theta)$  for random variable  $X$  and  $Z$ 
    - $p(x, z; \theta) = p(z; \theta)p(x|z; \theta)$
  - Dataset  $D = \{x^{(i)}\}$  for  $X$ , variable  $Z$  is never observed

- Maximal Likelihood Learning

$$L(\theta) = \log \prod_{x \in D} p(x; \theta) = \sum_{x \in D} \log \sum_z p(x, z; \theta)$$

- Marginal probability can be expensive to compute!
    - When  $z$  is continuous, the objective even becomes intractable
- Goal: a fast **approximation** of the **marginal probability**
  - Remark:  $L(\theta)$  is tractable when  $p(x, z) \propto \exp(-E(x, z))$

# Latent Variable Model: Training

- Goal: **approximation** of  $\log \sum_z p(x, z; \theta)$

- Idea#1: Importance Sampling

- Proposal distribution  $q(z)$

$$p(x) = \sum_z q(z) \cdot \frac{p(x, z; \theta)}{q(z)}$$


- The probability can be approximated by drawing samples from  $q(z)$

- Learning objective  $L(x; \theta)$

$$L(x; \theta) = \log \sum_z q(z) \cdot \frac{p(x, z; \theta)}{q(z)}$$


# Latent Variable Model: Training

- Goal: approximation of  $\log \sum_z q(z) \cdot \frac{p(x,z;\theta)}{q(z)}$
- Idea#2: concavity of  $\log(\cdot)$ 
  - $\log \sum_z q(z) \cdot \frac{p(x,z;\theta)}{q(z)}$
  - For any  $0 < x_1 \leq x_2 \leq 1$ ,
    - $\log(\alpha x_1 + (1 - \alpha)x_2) \geq \alpha \log(x_1) + (1 - \alpha) \log(x_2)$
  - More general, for any weights  $\alpha_i > 0$  &  $\sum_i \alpha_i = 1$ ,
    - $\log(\sum_i \alpha_i x_i) \geq \sum_i \alpha_i \log(x_i)$



# Latent Variable Model: Training

- Goal: approximation of  $\log \sum_z q(z) \cdot \frac{p(x,z;\theta)}{q(z)}$
- Idea#2: concavity of  $\log(\cdot)$ 
  - $\log \sum_z q(z) \cdot \frac{p(x,z;\theta)}{q(z)} \geq \sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)}$
  - For any  $0 < x_1 \leq x_2 \leq 1$ ,
    - $\log(\alpha x_1 + (1 - \alpha)x_2) \geq \alpha \log(x_1) + (1 - \alpha) \log(x_2)$
  - More general, for any weights  $\alpha_i > 0$  &  $\sum_i \alpha_i = 1$ ,
    - $\log(\sum_i \alpha_i x_i) \geq \sum_i \alpha_i \log(x_i)$





# Latent Variable Model: Training

- Goal: approximate  $\log \sum_z p(x, z; \theta)$ 
  - Ideas: importance sampling & concavity of  $\log(\cdot)$


- Evidence Lower Bound (ELBO)

$$\log p(x; \theta) = \log \sum_z p(x, z; \theta) \geq \sum_z q(z) \log \frac{p(x, z; \theta)}{q(z)}$$

- A tractable lower bound of the true objective
    - Easy to optimize
- When will the equality hold?
  - i.e., a tight lower bound
  - Sol:  $q(z) \leftarrow p(z|x; \theta)$


# Evidence Lower Bound

- ELBO becomes exact when  $q(z) = p(z|x; \theta)$ 
  - $\sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)} = \sum_z q(z) \log \frac{p(x,z;\theta)}{p(z|x;\theta)}$
  - $= \sum_z q(z) \log p(x; \theta)$
  - $= \log p(x; \theta)$
- We can optimize a tight lower bound by setting  $q(z) = p(z|x; \theta)$
- An iterative process
  - Optimize  $p(x, z; \theta)$  w.r.t. fixed  $q(z)$ 
    - $J(\theta) = \sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)}$
  - Set  $q(z) \leftarrow p(z|x; \theta)$
  - Repeat




# Evidence Lower Bound

- ELBO becomes exact when  $q(z) = p(z|x; \theta)$ 
  - $\sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)} = \sum_z q(z) \log \frac{p(x,z;\theta)}{p(z|x;\theta)}$
  - $= \sum_z q(z) \log p(x; \theta)$
  - $= \log p(x; \theta)$
- We can optimize a tight lower bound by setting  $q(z) = p(z|x; \theta)$
- An iterative process
  - Optimize  $p(x, z; \theta)$  w.r.t. fixed  $q(z)$ 
    - $J(\theta) = \sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)}$
  - Set  $q(z) \leftarrow p(z|x; \theta^0)$
  - Repeat




# Evidence Lower Bound

- ELBO becomes exact when  $q(z) = p(z|x; \theta)$ 
  - $\sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)} = \sum_z q(z) \log \frac{p(x,z;\theta)}{p(z|x;\theta)}$
  - $= \sum_z q(z) \log p(x; \theta)$
  - $= \log p(x; \theta)$
- We can optimize a tight lower bound by setting  $q(z) = p(z|x; \theta)$
- An iterative process
  - **Optimize  $p(x, z; \theta)$  w.r.t. fixed  $q(z; \theta^0)$** 
    - $J(\theta) = \sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)}$
  - Set  $q(z) \leftarrow p(z|x; \theta)$
  - Repeat




# Evidence Lower Bound

- ELBO becomes exact when  $q(z) = p(z|x; \theta)$ 
  - $\sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)} = \sum_z q(z) \log \frac{p(x,z;\theta)}{p(z|x;\theta)}$
  - $= \sum_z q(z) \log p(x; \theta)$
  - $= \log p(x; \theta)$
- We can optimize a tight lower bound by setting  $q(z) = p(z|x; \theta)$
- An iterative process
  - Optimize  $p(x, z; \theta)$  w.r.t. fixed  $q(z)$ 
    - $J(\theta) = \sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)}$
  - Set  $q(z) \leftarrow p(z|x; \theta^1)$
  - Repeat




# Evidence Lower Bound

- ELBO becomes exact when  $q(z) = p(z|x; \theta)$ 
  - $\sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)} = \sum_z q(z) \log \frac{p(x,z;\theta)}{p(z|x;\theta)}$
  - $= \sum_z q(z) \log p(x; \theta)$
  - $= \log p(x; \theta)$
- We can optimize a tight lower bound by setting  $q(z) = p(z|x; \theta)$
- An iterative process
  - Optimize  $p(x, z; \theta)$  w.r.t. fixed  $q(z; \theta^1)$ 
    - $J(\theta) = \sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)}$
  - Set  $q(z) \leftarrow p(z|x; \theta^2)$
  - Repeat




# Evidence Lower Bound

- ELBO becomes exact when  $q(z) = p(z|x; \theta)$ 
  - $\sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)} = \sum_z q(z) \log \frac{p(x,z;\theta)}{p(z|x;\theta)}$
  - $= \sum_z q(z) \log p(x; \theta)$
  - $= \log p(x; \theta)$
- We can optimize a tight lower bound by setting  $q(z) = p(z|x; \theta)$
- An iterative process
  - Optimize  $p(x, z; \theta)$  w.r.t. fixed  $q(z)$ 
    - $J(\theta) = \sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)}$
  - Set  $q(z) \leftarrow p(z|x; \theta)$
  - Repeat
  - **Converge to a local optimum**



# Evidence Lower Bound

- ELBO becomes exact when  $q(z) = p(z|x; \theta)$ 
  - $\sum_z q(z) \log \frac{p(x,z;\theta)}{q(z)} = \sum_z q(z) \log \frac{p(x,z;\theta)}{p(z|x;\theta)}$
  - $= \sum_z q(z) \log p(x; \theta)$
  - $= \log p(x; \theta)$
- We can optimize a tight lower bound by setting  $q(z) = p(z|x; \theta)$
- An iterative process
  - Optimize  $p(x, z; \theta)$  w.r.t. fixed  $q(z)$  (M-step)
  - Set  $q(z) \leftarrow p(z|x; \theta)$  (E-step)
  - **An EM algorithm**
- How to set  $q(z) \leftarrow p(z|x; \theta)$ ?





# Variational Inference

- Goal:  $q(z; \phi) \leftarrow p(z|x)$ 
  - Find a parameterized distribution  $q(z; \phi)$  to approximate the true posterior
    - In our case, approximate  $p(z|x; \theta)$  w.r.t. a fixed  $\theta$
  - Distance metric between  $q(z; \phi)$  and  $p(z|x)$ 
    - $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
- Variational Inference:  $\min_{\phi} KL(q||p)$

# Variational Inference

- Goal:  $q(z; \phi) \leftarrow p(z|x)$ 
  - Find a parameterized distribution  $q(z; \phi)$  to approximate the posterior
    - In our case, approximate  $p(z|x; \theta)$  w.r.t. a fixed  $\theta$
  - Distance metric between  $q(z; \phi)$  and  $p(z|x)$ 
    - $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
- Variational Inference:  $\min_{\phi} KL(q||p)$ 
  - Remark: pay attention to the order of KL (**reverse KL**)!
  - Mean-field variational inference
    - A factored proposal:  $q(z) = \prod_i q_i(z_i|x)$
    - By calculus of variation (变分法, 泛函分析领域)
$$\log q_i^*(z_i|x) = E_{z_{j \neq i}} [\log p(z, x)] + \text{constant}$$
    - Repeatedly update the distribution of  $q_i(z_i)$  using the expectation of  $p(z, x)$

# Variational Inference

- Goal:  $q(z; \phi) \leftarrow p(z|x)$ 
  - Find a parameterized distribution  $q(z; \phi)$  to approximate the posterior
    - In our case, approximate  $p(z|x; \theta)$  w.r.t. a fixed  $\theta$
  - Distance metric between  $q(z; \phi)$  and  $p(z|x)$ 
    - $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
- Variational Inference:  $\min_{\phi} KL(q||p)$ 
  - $KL(q(z; \phi)||p(z|x)) = \sum_z q(z; \phi) \log \frac{q(z; \phi)}{p(z|x)}$

# Variational Inference

- Goal:  $q(z; \phi) \leftarrow p(z|x)$ 
  - Find a parameterized distribution  $q(z; \phi)$  to approximate the posterior
    - In our case, approximate  $p(z|x; \theta)$  w.r.t. a fixed  $\theta$
  - Distance metric between  $q(z; \phi)$  and  $p(z|x)$ 
    - $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
- Variational Inference:  $\min_{\phi} KL(q||p)$ 
  - $KL(q(z; \phi)||p(z|x)) = \sum_z q(z; \phi) \log \frac{q(z; \phi)}{p(z|x)} = \sum_z q(z; \phi) \log \frac{q(z; \phi) \mathbf{p(x)}}{\mathbf{p(z,x)}}$

# Variational Inference

- Goal:  $q(z; \phi) \leftarrow p(z|x)$ 
  - Find a parameterized distribution  $q(z; \phi)$  to approximate the posterior
    - In our case, approximate  $p(z|x; \theta)$  w.r.t. a fixed  $\theta$
  - Distance metric between  $q(z; \phi)$  and  $p(z|x)$ 
    - $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
- Variational Inference:  $\min_{\phi} KL(q||p)$ 
  - $KL(q(z; \phi)||p(z|x)) = \sum_z q(z; \phi) \log \frac{q(z; \phi)}{p(z|x)} = \sum_z q(z; \phi) \log \frac{q(z; \phi) \mathbf{p(x)}}{p(z,x)}$
  - $= \sum_z q(z; \phi) \log \mathbf{p(x)} - \sum_z q(z; \phi) \log \frac{p(z,x)}{q(z; \phi)}$

# Variational Inference

- Goal:  $q(z; \phi) \leftarrow p(z|x)$ 
  - Find a parameterized distribution  $q(z; \phi)$  to approximate the posterior
    - In our case, approximate  $p(z|x; \theta)$  w.r.t. a fixed  $\theta$
  - Distance metric between  $q(z; \phi)$  and  $p(z|x)$ 
    - $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
- Variational Inference:  $\min_{\phi} KL(q||p)$ 
  - $KL(q(z; \phi)||p(z|x)) = \sum_z q(z; \phi) \log \frac{q(z; \phi)}{p(z|x)} = \sum_z q(z; \phi) \log \frac{q(z; \phi)p(x)}{p(z, x)}$
  - $= \text{Constant} \log p(x) - \sum_z q(z; \phi) \log \frac{p(z, x)}{q(z; \phi)}$

# Variational Inference

- Goal:  $q(z; \phi) \leftarrow p(z|x)$ 
    - Find a parameterized distribution  $q(z; \phi)$  to approximate the posterior
      - In our case, approximate  $p(z|x; \theta)$  w.r.t. a fixed  $\theta$
    - Distance metric between  $q(z; \phi)$  and  $p(z|x)$ 
      - $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
  - Variational Inference:  $\min_{\phi} KL(q||p)$ 
    - $KL(q(z; \phi)||p(z|x)) = \sum_z q(z; \phi) \log \frac{q(z; \phi)}{p(z|x)} = \sum_z q(z; \phi) \log \frac{q(z; \phi)p(x)}{p(z, x)}$
    - $= \log p(x) - \sum_z q(z; \phi) \log \frac{p(z, x)}{q(z; \phi)}$
    - $L(\phi) = \sum_z q(z; \phi) \log \frac{p(z, x)}{q(z; \phi)}$
- Evidence Lower Bound (ELBO)!!!**
- Also called variational lower bound in VI*

# Variational Inference

- Goal:  $q(z; \phi) \leftarrow p(z|x)$ 
  - Find a parameterized distribution  $q(z; \phi)$  to approximate the posterior
    - In our case, approximate  $p(z|x; \theta)$  w.r.t. a fixed  $\theta$
  - Distance metric between  $q(z; \phi)$  and  $p(z|x)$ 
    - $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
- Variational Inference:  $\min_{\phi} KL(q||p)$ 
  - $KL(q(z; \phi)||p(z|x)) = \sum_z q(z; \phi) \log \frac{q(z; \phi)}{p(z|x)} = \sum_z q(z; \phi) \log \frac{q(z; \phi)p(x)}{p(z, x)}$
  - $= \log p(x) - \sum_z q(z; \phi) \log \frac{p(z, x)}{q(z; \phi)}$
  - $L(\phi) = \sum_z q(z; \phi) \log \frac{p(z, x)}{q(z; \phi)}$



# Variational Inference

- Goal:  $q(z; \phi) \leftarrow p(z|x)$ 
  - Find a parameterized distribution  $q(z; \phi)$  to approximate the posterior
    - In our case, approximate  $p(z|x; \theta)$  w.r.t. a fixed  $\theta$
  - Distance metric between  $q(z; \phi)$  and  $p(z|x)$ 
    - $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
- Variational Inference:  $\min_{\phi} KL(q||p)$ 
  - $\log p(x) = KL(q(z; \phi)||p(z|x)) + \sum_z q(z; \phi) \log \frac{p(z,x)}{q(z;\phi)}$
  - $\quad = \text{approximate error} + \text{ELBO} \geq \text{ELBO}$
  - $L(\phi) = \sum_z q(z; \phi) \log \frac{p(z,x)}{q(z;\phi)}$

# Variational Inference (Explained)

- General Formulation of Bayesian Inference
  - Dataset  $D = \{x\}$
  - Model  $p(x; \theta)$  with parameter  $\theta$
  - Goal  $p(\theta|x)$ 
    - Remark: optimization learns a single  $\theta^*$  while BI learns a distribution
- Variational Inference as a Mean of Approximate Bayesian Inference
  - Use  $q(\theta; \phi)$  to approximate  $p(\theta|x)$
  - VI Objective:  $KL(q||p) = C + ELBO$
  - Interpretation: VI objective is a *lower bound* of  $\log p(x)$   
 $\log p(x; \theta) = \text{approximation error} + ELBO$
- VAE naturally inherits all the nice mathematical properties of VI 😊
  - Further read: black-box variational inference <https://arxiv.org/abs/1401.0118>

# Latent Variable Model: Training

- Latent Variable Model:  $p(z, x) = p(z)p(x|z)$ 
  - MLE objective:  $p(x; \theta) = \sum_z p(z, x; \theta)$
- ELBO:  $p(x; \theta) \geq \sum_z q(z) \log \frac{p(x, z; \theta)}{q(z)} = L(\theta; q)$ 
  - Iterative learning: (1) optimize  $\theta$  and (2)  $q(z) \leftarrow p(z|x; \theta)$

# Latent Variable Model: Training

- Latent Variable Model:  $p(z, x) = p(z)p(x|z)$ 
  - MLE objective:  $p(x; \theta) = \sum_z p(z, x; \theta)$
- ELBO:  $p(x; \theta) \geq \sum_z q(z) \log \frac{p(z, x; \theta)}{q(z)} = L(\theta; q)$ 
  - Iterative learning: (1) optimize  $\theta$  and (2)  $q(z) \leftarrow p(z|x; \theta)$
- Variational Inference
  - Approximate  $p(z|x; \theta)$  by a tractable distribution  $q(z; \phi)$

# Latent Variable Model: Training

- Latent Variable Model:  $p(z, x) = p(z)p(x|z)$ 
  - MLE objective:  $p(x; \theta) = \sum_z p(z, x; \theta)$
- ELBO:  $p(x; \theta) \geq \sum_z q(z) \log \frac{p(z, x; \theta)}{q(z)} = L(\theta; q)$ 
  - Iterative learning: (1) optimize  $\theta$  and (2)  $q(z) \leftarrow p(z|x; \theta)$
- Variational Inference
  - Approximate  $p(z|x; \theta)$  by a tractable distribution  $q(z; \phi)$

$$L(\phi; \theta) = \sum_z q(z; \phi) \log \frac{p(z, x; \theta)}{q(z; \phi)}$$

# Latent Variable Model: Training

- Latent Variable Model:  $p(z, x) = p(z)p(x|z)$ 
  - MLE objective:  $p(x; \theta) = \sum_z p(z, x; \theta)$
- ELBO:  $p(x; \theta) \geq \sum_z q(z) \log \frac{p(z, x; \theta)}{q(z)} = L(\theta; q)$ 
  - Iterative learning: (1) optimize  $\theta$  and (2)  $q(z) \leftarrow p(z|x; \theta)$
- Variational Inference
  - Approximate  $p(z|x; \theta)$  by a tractable distribution  $q(z; \phi)$ 
$$L(\phi; \theta) = \sum_z q(z; \phi) \log \frac{p(z, x; \theta)}{q(z; \phi)}$$

**Use VI to learn a separate  $q(z; \phi)$  for each possible  $x$ ?**

# Latent Variable Model: Training

- Latent Variable Model:  $p(z, x) = p(z)p(x|z)$ 
  - MLE objective:  $p(x; \theta) = \sum_z p(z, x; \theta)$
- ELBO:  $p(x; \theta) \geq \sum_z q(z) \log \frac{p(z, x; \theta)}{q(z)} = L(\theta; q)$ 
  - Iterative learning: (1) optimize  $\theta$  and (2)  $q(z) \leftarrow p(z|x; \theta)$
- **Amortized** Variational Inference
  - Approximate  $p(z|x; \theta)$  by a **conditional** tractable distribution  $q(z|x; \phi)$ 
$$L(\phi; \theta) = \sum_z q(z|x; \phi) \log \frac{p(z, x; \theta)}{q(z|x; \phi)}$$
  - A universal approximation  $q$  for any  $x$  and  $p(z|x)$

# Latent Variable Model: Training

- Latent Variable Model:  $p(z, x) = p(z)p(x|z)$ 
  - MLE objective:  $p(x; \theta) = \sum_z p(z, x; \theta)$
- ELBO:  $p(x; \theta) \geq \sum_z q(z) \log \frac{p(x, z; \theta)}{q(z)} = L(\theta; q)$ 
  - Iterative learning: (1) optimize  $\theta$  and (2)  $q(z) \leftarrow p(z|x; \theta)$
- Amortized Variational Inference
  - Approximate  $p(z|x; \theta)$  by a conditional tractable distribution  $q(z|x; \phi)$

$$L(\phi; \theta) = \sum_z q(z|x; \phi) \log \frac{p(z, x; \theta)}{q(z|x; \phi)}$$

- Joint Learning  $J(\theta, \phi; x)$

$$J(\theta, \phi; x) = \sum_z q(z|x; \phi) \log \frac{p(z, x; \theta)}{q(z|x; \phi)}$$



# Latent Variable Model: Training

- Learning objective  $J(\theta, \phi; x)$ 
  - $J(\theta, \phi; x) = \sum_z q(z|x; \phi)(\log p(z, x; \theta) - \log q(z|x; \phi))$

# Latent Variable Model: Training

- Learning objective  $J(\theta, \phi; x)$ 
  - $J(\theta, \phi; x) = \sum_z q(z|x; \phi)(\log p(z, x; \theta) - \log q(z|x; \phi))$
  - $\quad \quad \quad = \sum_z q(z|x; \phi)(\log p(x|z; \theta) - \log q(z|x; \phi) + \log p(z; \theta))$

# Latent Variable Model: Training

- Learning objective  $J(\theta, \phi; x)$ 
  - $J(\theta, \phi; x) = \sum_z q(z|x; \phi)(\log p(z, x; \theta) - \log q(z|x; \phi))$
  - $= \sum_z q(z|x; \phi)(\log p(x|z; \theta) - \log q(z|x; \phi) + \log p(z; \theta))$
  - $= \sum_z q(z|x; \phi) \log p(x|z; \theta) - \sum_z q(z|x; \phi) \log \frac{q(z|x; \phi)}{p(z; \theta)}$

# Latent Variable Model: Training

- Learning objective  $J(\theta, \phi; x)$ 
  - $J(\theta, \phi; x) = \sum_z q(z|x; \phi)(\log p(z, x; \theta) - \log q(z|x; \phi))$
  - $= \sum_z q(z|x; \phi)(\log p(x|z; \theta) - \log q(z|x; \phi) + \log p(z; \theta))$
  - $= \sum_z q(z|x; \phi) \log p(x|z; \theta) - \sum_z q(z|x; \phi) \log \frac{q(z|x; \phi)}{p(z; \theta)}$
  - $= \underbrace{E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)]}_{\text{Expectation of log likelihood (reconstruction)}} - \underbrace{KL(q(z|x; \phi) || p(z; \theta))}_{\text{KL divergence}}$
- Design of  $p(z, x; \theta)$  and  $q(z|x; \phi)$ 
  - Principle: easy to compute!
  - Gaussian prior:  $p(z) \sim N(0, I)$
  - Gaussian likelihood:  $p(x_{ij}|z; \theta) \sim N(f_{ij}(z; \theta), 1)$
  - Isomorphic Gaussian:  $q(z|x; \phi) \sim N(\mu(x; \phi), \text{diag}(\exp(\sigma(x; \phi))))$

# Latent Variable Model: Training

- Learning objective  $J(\theta, \phi; x)$

- $J(\theta, \phi; x) = \sum_z q(z|x; \phi)(\log p(z, x; \theta) - \log q(z|x; \phi))$
  - $= \sum_z q(z|x; \phi)(\log p(x|z; \theta) - \log q(z|x; \phi) + \log p(z; \theta))$
  - $= \sum_z q(z|x; \phi) \log p(x|z; \theta) - \sum_z q(z|x; \phi) \log \frac{q(z|x; \phi)}{p(z; \theta)}$
  - $= E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)] - KL(q(z|x; \phi) || p(z; \theta))$

Expectation of log likelihood (reconstruction)

KL divergence

- Design of  $p(z, x; \theta)$  and  $q(z|x; \phi)$

- Principle: easy to compute!
  - Gaussian prior:  $p(z) \sim N(0, I)$
  - Gaussian likelihood:  $p(x_{ij}|z; \theta) \sim N(f_{ij}(z; \theta), 1)$
  - Isomorphic Gaussian:  $q(z|x; \phi) \sim N(\mu(x; \phi), \text{diag}(\exp(\sigma(x; \phi))))$

**Neural networks!**


# Variational Autoencoder

- VAE Architecture

- Isomorphic Gaussian:  $q(z|x; \phi) \sim N(\mu(x; \phi), \text{diag}(\exp(\sigma(x; \phi))))$
- Gaussian prior:  $p(z) \sim N(0, I)$
- Gaussian likelihood:  $p(x|z; \theta) \sim N(f(z; \theta), I)$

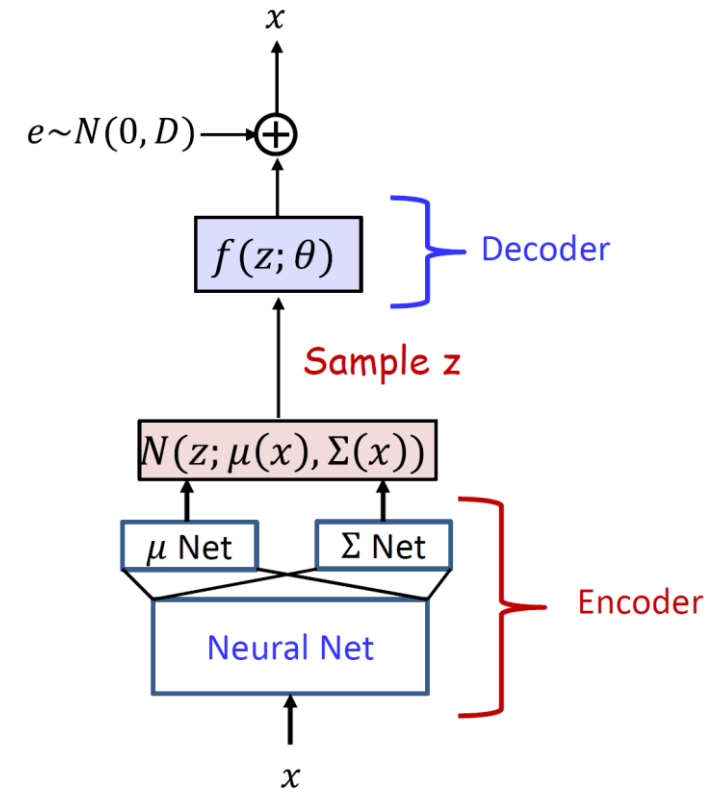
- Autoencoder  $x \rightarrow z \rightarrow x$

- Unsupervised learning (data to data,  $z$  never observed)
- Encoder  $q(z|x; \phi): x \rightarrow z$
- Decoder  $p(x|z; \theta): z \rightarrow x$
- Remark
  - $p(x|z)$  is the actual generative model
  - $q(z|x)$  is only the proposal
    - but optimized to approximate  $p(z|x)$



# Variational Autoencoder

- Training via jointly optimizing ELBO
  - $J(\phi, \theta; x) = \mathbb{E}_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)] - KL(q(z|x; \phi) || p(z))$
  - Two terms: likelihood term & KL term



# Variational Autoencoder

- Training via jointly optimizing ELBO


- $J(\phi, \theta; x) = \mathbb{E}_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)] - \textcolor{red}{KL}(q(z|x; \phi) || p(z))$

- KL penalty

- $q(z|x; \phi) \sim N\left(\mu(x; \phi), \text{diag}(\exp(\sigma(x; \phi)))\right)$
  - $p(z) \sim N(0, I)$
  - Closed-form!

$$D_{\text{KL}}(\mathcal{N}_0 || \mathcal{N}_1) = \frac{1}{2} \left\{ \text{tr}(\mathbf{\Sigma}_1^{-1} \mathbf{\Sigma}_0) + (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_0)^T \mathbf{\Sigma}_1^{-1} (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_0) - k + \ln \frac{|\mathbf{\Sigma}_1|}{|\mathbf{\Sigma}_0|} \right\}$$

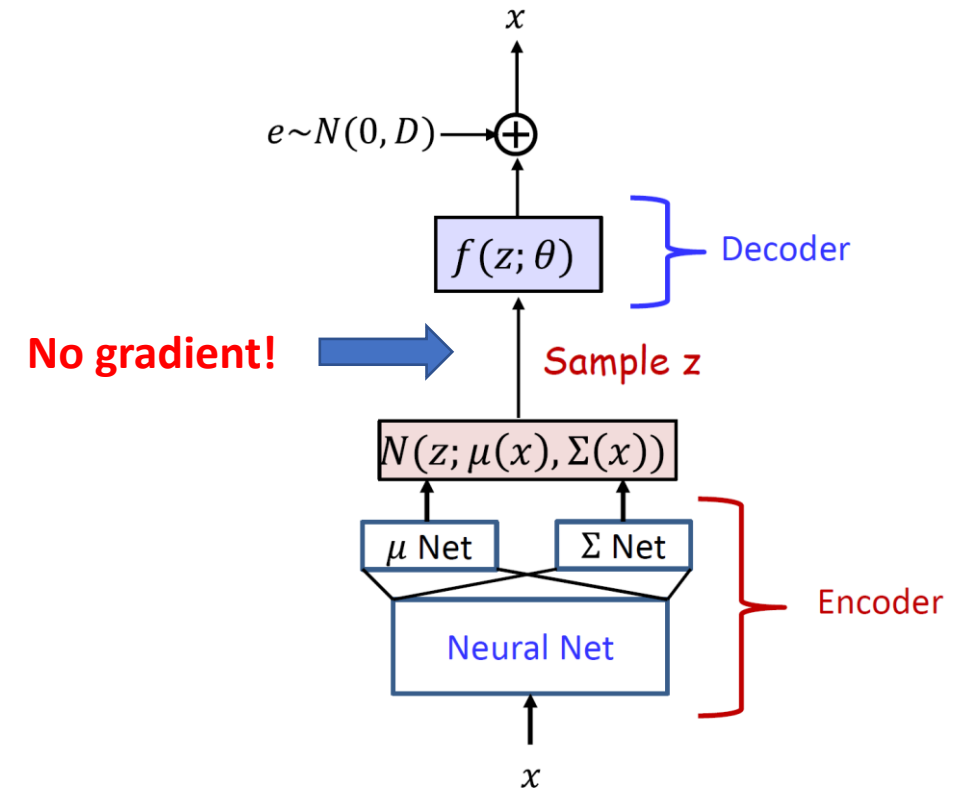
- Implement it in your coding project 😊






# Variational Autoencoder

- Training via jointly optimizing ELBO
  - $J(\phi, \theta; x) = \mathbb{E}_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)] - KL(q(z|x; \phi) || p(z))$
- Likelihood term (reconstruction loss)
  - Monte-Carlo estimate!
    - Draw samples from  $q(z|x; \phi)$
    - Compute gradient of  $\theta$ :  $L(\theta) \propto \sum_z |x - f(z; \theta)|^2$ 
      - $x \sim N(f(z; \theta); I)$
      - $p(x) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2} |x - f(z; \theta)|^2\right)$
    - How to get the gradient of  $\phi$  through  $q(z; \phi)$ ??  
 $L(\phi) = \mathbb{E}_{z \sim q(z; \phi)} [\log p(x|z)]$



# Variational Autoencoder

- Training via jointly optimizing ELBO
  - $J(\phi, \theta; x) = \mathbb{E}_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)] - KL(q(z|x; \phi) || p(z))$
- Likelihood term (reconstruction loss)
  - Monte-Carlo estimate!
    - Draw samples from  $q(z|x; \phi)$
    - Compute gradient of  $\theta: L(\theta) \propto \sum_z |x - f(z; \theta)|^2$
  - Re-parameterization trick
    - Recap in autoregressive flow
      - $z \sim N(\mu, \sigma^2) \iff z = \mu + \sigma \cdot \epsilon, \epsilon \sim N(0,1)$
    - $L(\phi) \propto \mathbb{E}_{z \sim q(z|\phi)} [|f(z) - x|^2]$
    - $\propto \mathbb{E}_{\epsilon \sim N(0, I)} [|f(\mu(x; \phi) + \sigma(x; \phi) \cdot \epsilon) - x|^2]$ 
      - Monte-Carlo estimate for  $\nabla L(\phi)$ !
      - 1 sample for  $\epsilon$  is sufficient for stable training



# Variational Autoencoder

- Variational Autoencoder (VAE)
  - Encoder  $q(z|x; \phi)$
  - Decoder  $p(x|z; \theta)$
  - End-to-end unsupervised learning ( $x \rightarrow z \sim q(z|x) \rightarrow x$ )  
 $J(\phi, \theta; x) = \mathbb{E}_{\epsilon \sim N(0, I)} [\log p(x|\mu(x; \phi) + \sigma(x; \phi) \cdot \epsilon; \theta)] - KL(q(z; \phi) || p(z))$
  - By Kingma & Welling, ICLR 2013 (34k citation)

---

## Auto-Encoding Variational Bayes

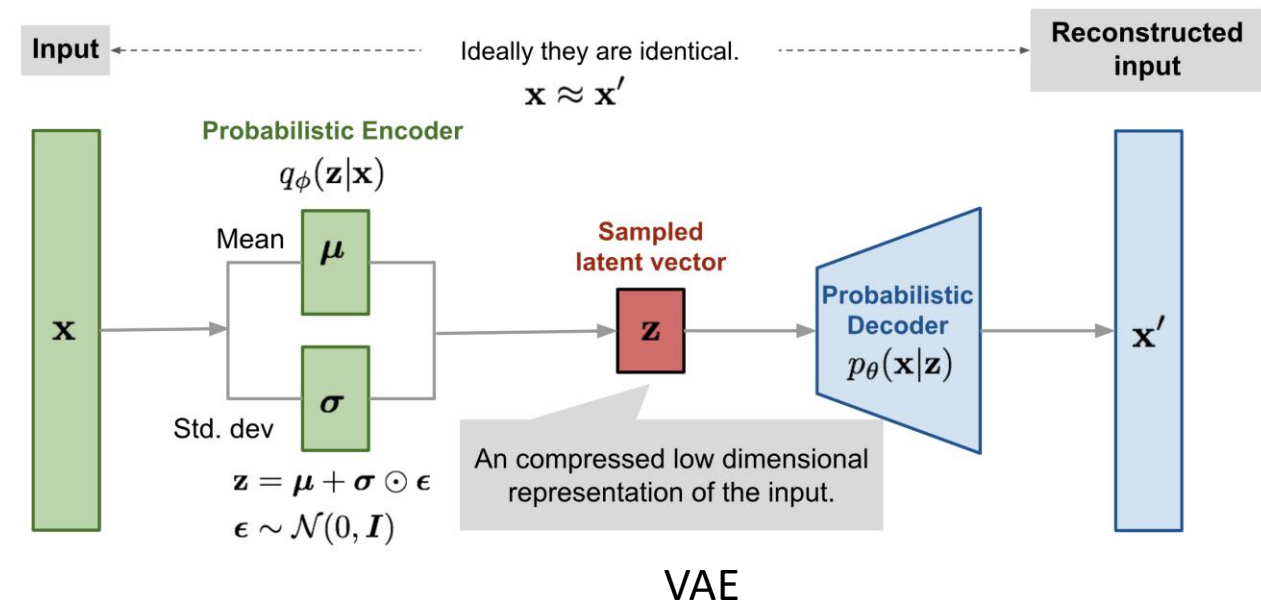
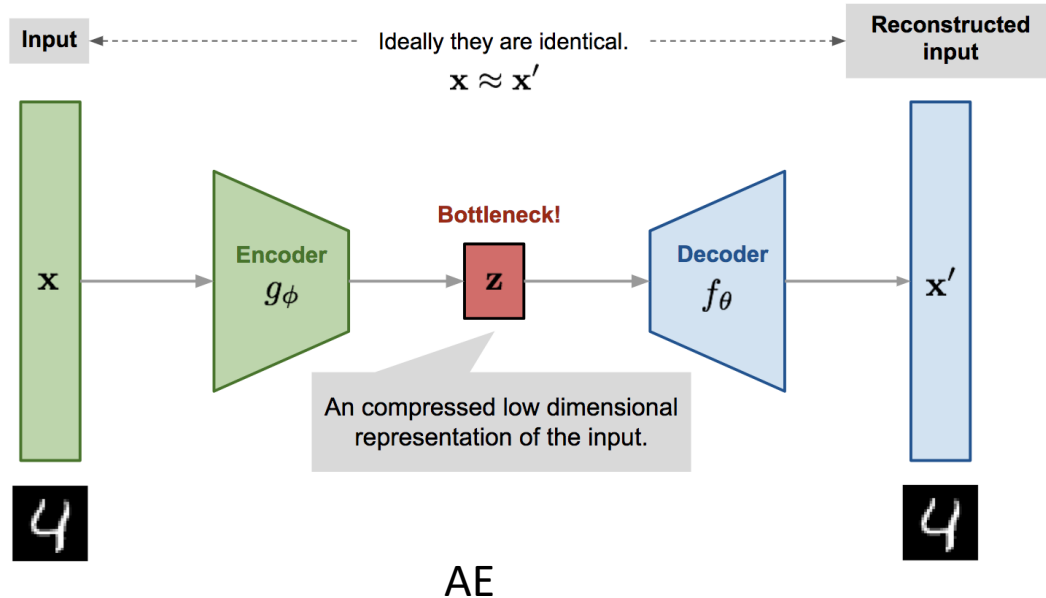
---

Diederik P. Kingma  
Machine Learning Group  
Universiteit van Amsterdam  
dpkingma@gmail.com

Max Welling  
Machine Learning Group  
Universiteit van Amsterdam  
welling.max@gmail.com

# VAE v.s. Standard AE

- Autoencoder
  - A classical unsupervised learning method for representation learning
- VAE: a simple generative extension of AE
  - Generative model: AE + Gaussian noise on  $z$
  - KL penalty: L2 constraint on the latent vector  $z$

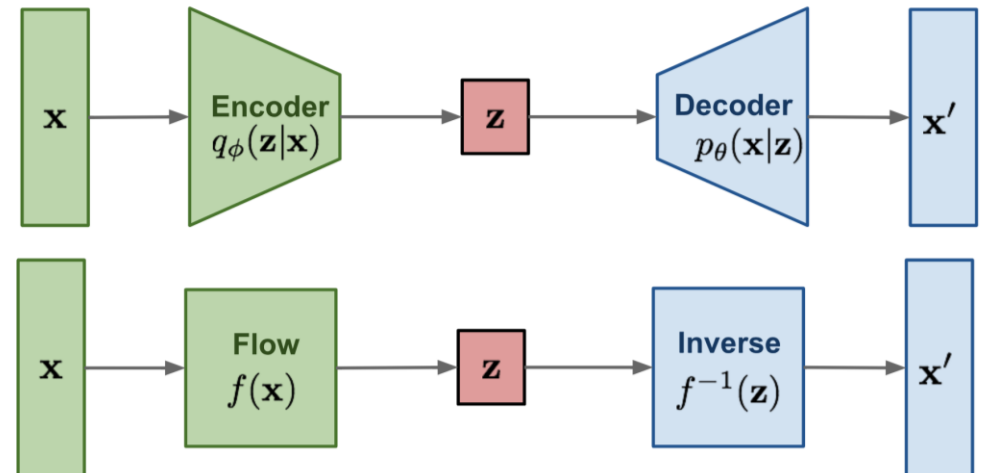


# VAE v.s. Flow Model

- Both model has a latent representation  $z$
- Flow model
  - Encoder: inference mapping; decoder: generation mapping
  - Exact inference but no dimension reduction!
- VAE
  - Approximate inference
  - Dimension reduction
  - Flexible architecture



VAE: maximize ELBO.

Flow-based  
generative models:  
minimize the negative  
log-likelihood



# Variational Autoencoder

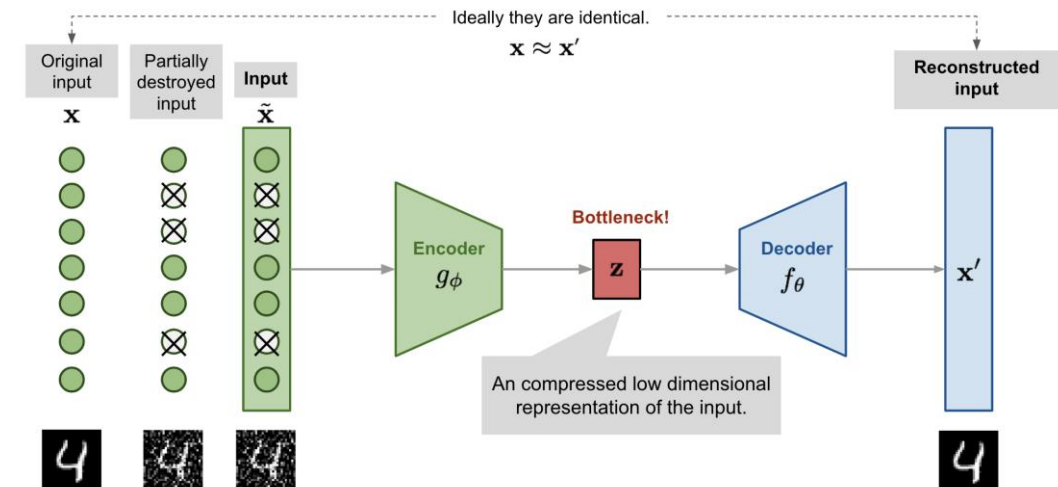

- Interpretable latent space
  - By interpolating  $z$ , we can observe how the generated samples vary
  - Automatic clustering in the (low-dimensional) latent space







# Inpainting with VAE

- Inference the missing pixels?
  - Fully observable training data  $D = \{x^{(i)}\}$
  - Standard VAE:  $q(z|x)$  &  $p(x|z)$
  - Corrupted data:  $\bar{x} = x \odot \text{mask}$
  - Goal:  $q(z|\bar{x}) \approx q(z|x)$ 
    - We do not need to change the generator  $p(x|z)$
- Randomized mask in training!
  - $x \odot \text{mask} \rightarrow z \rightarrow x$
  - Better encoder architecture
    - Masked convolution
    - Idea: convolution only on unmasked pixels
    - Image Inpainting for Irregular Holes Using Partial Convolutions (ECCV 2018)



# Conditioned VAE

- Include label in VAE
  - $D = \{(x^{(i)}, y^{(i)})\}$
  - Encoder:  $q(z|x, y; \phi)$
  - Decoder:  $p(x|y, z; \theta)$
  - Conditioned generation!



- What if we have both labeled data and unlabeled data?




# Conditioned VAE

- Semi-supervised learning


- $D_l = \{(x^{(i)}, y^{(i)})\}$
- $D_u = \{(x^{(i)})\}$
- Decoder:  $p(x|y, z; \theta)$
- Encoder?

- $q(z, y|x; \phi)$

- In practice:  $q(z, y|x; \phi) = q(z|x; \phi) \cdot q(y|x; \phi)$



conditioned VAE



conditioned generation



# Conditioned VAE

- Semi-supervised learning

- $D_l = \{(x^{(i)}, y^{(i)})\}$
- $D_u = \{(x^{(i)})\}$
- Decoder:  $p(x|y, z; \theta)$
- Encoder:  $q(z, y|x; \phi)$

- Training

- Easy on supervised data
  - Cross-entropy loss on  $q(y|x; \phi)$  on labeled data
- What about unlabeled data?





# Conditioned VAE

- Semi-supervised learning

- $D_l = \{(x^{(i)}, y^{(i)})\}$
- $D_u = \{(x^{(i)})\}$
- Decoder:  $p(x|y, z; \theta)$
- Encoder:  $q(z, y|x; \phi)$

- Training on  $D_u$

- Loss = reconstruction + KL penalty
- KL penalty:  $KL(q(z)||p(z)) + KL(q(y)||p(y))$  ( $p(y) \sim \text{uniform}$ )
- Reconstruction loss:  $L = E_{z, y \sim q(z, y)} [\log p(x|z, y; \theta)]$




# Conditioned VAE

- Semi-supervised learning


- $D_l = \{(x^{(i)}, y^{(i)})\}$
- $D_u = \{(x^{(i)})\}$
- Decoder:  $p(x|y, z; \theta)$
- Encoder:  $q(z, y|x; \phi)$

- Training on  $D_u$

- Loss = reconstruction + KL penalty
- KL penalty:  $KL(q(z)||p(z)) + KL(q(y)||p(y))$  ( $p(y) \sim \text{uniform}$ )
- Reconstruction loss:  $L = E_{\epsilon \sim N(0, I), y \sim q(y)} [\log p(x|\mu(x) + \sigma(x) \cdot \epsilon, y; \theta)]$ 
  - Reparameterization trick for  $z$



conditioned VAE




conditioned generation


# Conditioned VAE

- Semi-supervised learning

- $D_l = \{(x^{(i)}, y^{(i)})\}$
- $D_u = \{(x^{(i)})\}$
- Decoder:  $p(x|y, z; \theta)$
- Encoder:  $q(z, y|x; \phi)$



conditioned VAE



conditioned generation


- Training on  $D_u$

- Loss = reconstruction + KL penalty
- KL penalty:  $KL(q(z)||p(z)) + KL(q(y)||p(y))$  ( $p(y) \sim \text{uniform}$ )
- Reconstruction loss:  $L = E_{\epsilon \sim N(0, I), y \sim q(y)} [\log p(x|\mu(x) + \sigma(x) \cdot \epsilon, y; \theta)]$ 
  - Reparameterization trick for  $z$
  - What about  $y$  ?
    - .... although we do have tricks in lecture 11 :P


# Conditioned VAE

- Semi-supervised learning

- $D_l = \{(x^{(i)}, y^{(i)})\}$
- $D_u = \{(x^{(i)})\}$
- Decoder:  $p(x|y, z; \theta)$
- Encoder:  $q(z, y|x; \phi)$



conditioned VAE



conditioned generation


- Training on  $D_u$

- Loss = reconstruction + KL penalty
- KL penalty:  $KL(q(z)||p(z)) + KL(q(y)||p(y))$  ( $p(y) \sim \text{uniform}$ )
- Reconstruction loss:  $L = E_{\epsilon \sim N(0, I), y \sim q(y)} [\log p(x|\mu(x) + \sigma(x) \cdot \epsilon, y; \theta)]$ 
  - Reparameterization trick for  $z$
  - **We only have a few labels! Expand the expectation!**


# Conditioned VAE

- Semi-supervised learning

- $D_l = \{(x^{(i)}, y^{(i)})\}$
- $D_u = \{(x^{(i)})\}$
- Decoder:  $p(x|y, z; \theta)$
- Encoder:  $q(z, y|x; \phi)$



conditioned VAE



conditioned generation


- Training on  $D_u$

- Loss = reconstruction + KL penalty
- KL penalty:  $KL(q(z)||p(z)) + KL(q(y)||p(y))$  ( $p(y) \sim \text{uniform}$ )
- Reconstruction loss:
  - $L = E_{\epsilon \sim N(0, I), y \sim q(y)} [\log p(x|\mu(x) + \sigma(x) \cdot \epsilon, y; \theta)]$
  - $= E_{\epsilon \sim N(0, I)} [\sum_c q(y = c) \cdot \log p(x|\mu(x) + \sigma(x) \cdot \epsilon, y; \theta)]$


# Conditioned VAE

- Semi-supervised learning

- $D_l = \{(x^{(i)}, y^{(i)})\}$
- $D_u = \{(x^{(i)})\}$
- Decoder:  $p(x|y, z; \theta)$
- Encoder:  $q(z, y|x; \phi)$



conditioned VAE



conditioned generation

- Training on the entire dataset  $D$

- Supervised loss  $L^l$ 
  - Cross entropy for  $q(y)$ ; VAE loss for  $q(z)$  &  $p(x|z, y)$
- Unsupervised loss  $L^u$ 
  - Expanded likelihood over  $y$  for reconstruction loss
- Combined loss:  $J(\theta, \phi) = L^l + \beta L^u$ 
  - Leverage massive unlabeled data!

---

## Semi-supervised Learning with Deep Generative Models

---

Diederik P. Kingma\*, Danilo J. Rezende†, Shakir Mohamed†, Max Welling\*

\*Machine Learning Group, Univ. of Amsterdam, {D.P.Kingma, M.Welling}@uva.nl

†Google Deepmind, {danilor, shakir}@google.com




# Variational Autoencoder: Summary

- Pros
  - Flexible architecture & stable training
- Cons
  - Approximate inference


# Variational Autoencoder: Summary

- Pros
  - Flexible architecture & stable training
- Cons
  - **Approximate** inference
    - Intrinsic issue of KL divergence in VI
      - KL is asymmetric
        - VI:  $KL(q||p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$
      - $KL(q||p)$ : reverse (exclusive) KL
      - $KL(p||q)$ : forward (inclusive) KL
    - **The mode collapse issue**
      - Use forward KL?
      - Further reading of interest
      - <https://arxiv.org/abs/2202.01841>




(a)

inclusive KL



(b)


exclusive KL



(c)

# Variational Autoencoder:

- Pros
  - Flexible architecture & stable training
- Cons
  - **Approximate** inference
    - Intrinsic issue of KL divergence in VI
    - Assumed density of  $q(z|x)$  &  $p(z)$ 
      - $p(z) \sim N(0, I)$  for computation reason
        - We can have a more powerful prior (later in lecture 10)
        - E.g., structured VAE; VQ-VAE-2
      - $q(z|x) \sim N(\mu(x), \Sigma(x))$ 
        - What if  $p(z|x)$  is multi-modal?
        - We need a more powerful proposal distribution
          - E.g., flow models as  $q(z)$




(a) Predictions after 200 training epochs.

(b) Predictions after 1100 training epochs.

## Structured VAE

<https://arxiv.org/abs/1603.06277>



## VQ-VAE-2

<https://arxiv.org/abs/1906.00446>

---

## Variational Inference with Normalizing Flows

---

Danilo Jimenez Rezende  
Shakir Mohamed  
Google DeepMind, London

DANILOR@GOOGLE.COM  
SHAKIR@GOOGLE.COM


<https://arxiv.org/abs/1505.05770>

# Variational Autoencoder: Summary

- Pros
  - Flexible architecture & stable training
- Cons
  - **Approximate** inference
    - Intrinsic issue of KL divergence
    - Assumed density of  $q(z)$  &  $p(z)$
    - Variance due to single-step sampling
      - Importance-weighted autoencoder (Burda, Grosse & Ruslan, ICLR16)
        - <https://arxiv.org/abs/1509.00519>
      - Use more than one samples from  $q(z|x)$  for a tighter lower-bound


# Variational Autoencoder: Summary

- Pros
  - Flexible architecture & stable training
- Cons
  - **Approximate** inference
    - Intrinsic issue of KL divergence
    - Assumed density of  $q(z)$  &  $p(z)$
    - Variance due to single-step sampling
    - MLE as the reconstruction loss
      - $p(x|z; \theta) = N(f(z; \theta), I)$
      - Blurry samples!
        - Improve the decoder architecture
        - Balancing the KL penalty and reconstruction loss
        - Gaussian latent to discrete latent (in lecture 11)
        - Change the loss! (next lecture 😊)



# Variational Autoencoder: Summary

- Pros
  - Flexible architecture & stable training
- Cons
  - **Approximate** inference
    - Intrinsic issue of KL divergence
    - Assumed density of  $q(z)$  &  $p(z)$
    - Variance due to single-step sampling
    - MLE as the reconstruction loss
      - $p(x|z; \theta) = N(f(z; \theta), I)$
      - Blurry samples!
        - Improve the decoder architecture
        - **Balancing the KL penalty and reconstruction loss**
        - Gaussian latent to discrete latent (in lecture 10)
        - Change the loss! (next lecture 😊)



# VAE Variants

- VAE Objective (ELBO)

$$J(\theta, \phi; x) = \underbrace{E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)]}_{\text{Reconstruction}} - \underbrace{KL(q(z|x; \phi) || p(z; \theta))}_{\text{KL penalty}}$$



# VAE Variants


- $\beta$ -VAE (Higgins et. al, DeepMind, ICLR 2017)

$$J(\theta, \phi; x) = \underbrace{E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)]}_{\text{Reconstruction}} - \underbrace{\beta KL(q(z|x; \phi) || p(z; \theta))}_{\text{KL penalty}}$$

- Interpretation

$$\max_{\theta, \phi} E_{x \sim D} \left[ E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)] \right]$$

subject to  $KL(q(z|x; \phi) || p(z)) < \epsilon$





# VAE Variants

- $\beta$ -VAE (Higgins et. al, DeepMind, ICLR 2017)

$$J(\theta, \phi; x) = \underbrace{E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)]}_{\text{Reconstruction}} - \underbrace{\beta \text{KL}(q(z|x; \phi) || p(z; \theta))}_{\text{KL penalty}}$$

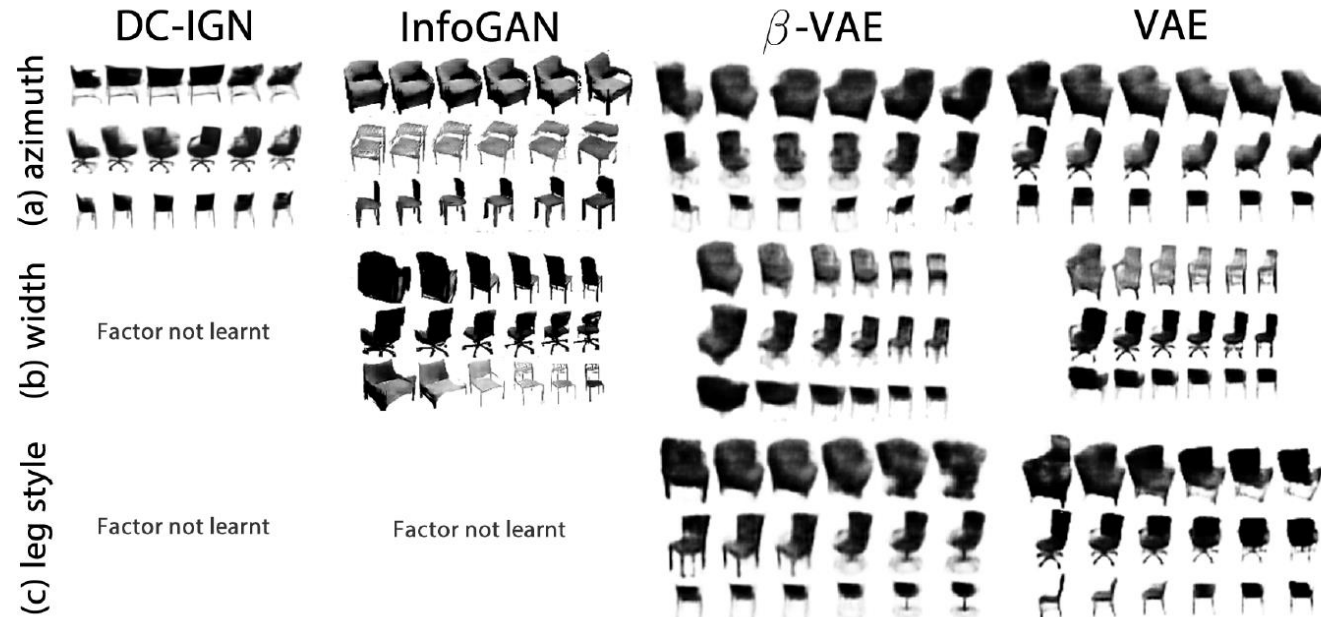
- Special cases
  - $\beta = 0$ : standard AE
  - $\beta = 1$ : standard VAE
  - $\beta > 1$ : force the latent space closer to isomorphic Gaussian
    - Insight: each dimension of  $z$  are forced to be independent
    - Disentangle factors!

# VAE Variants

- $\beta$ -VAE (Higgins et. al, DeepMind, ICLR 2017)

$$J(\theta, \phi; x) = \underbrace{E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)]}_{\text{Reconstruction}} - \underbrace{\beta KL(q(z|x; \phi) || p(z; \theta))}_{\text{KL penalty}}$$

- Learned factors in  $z$

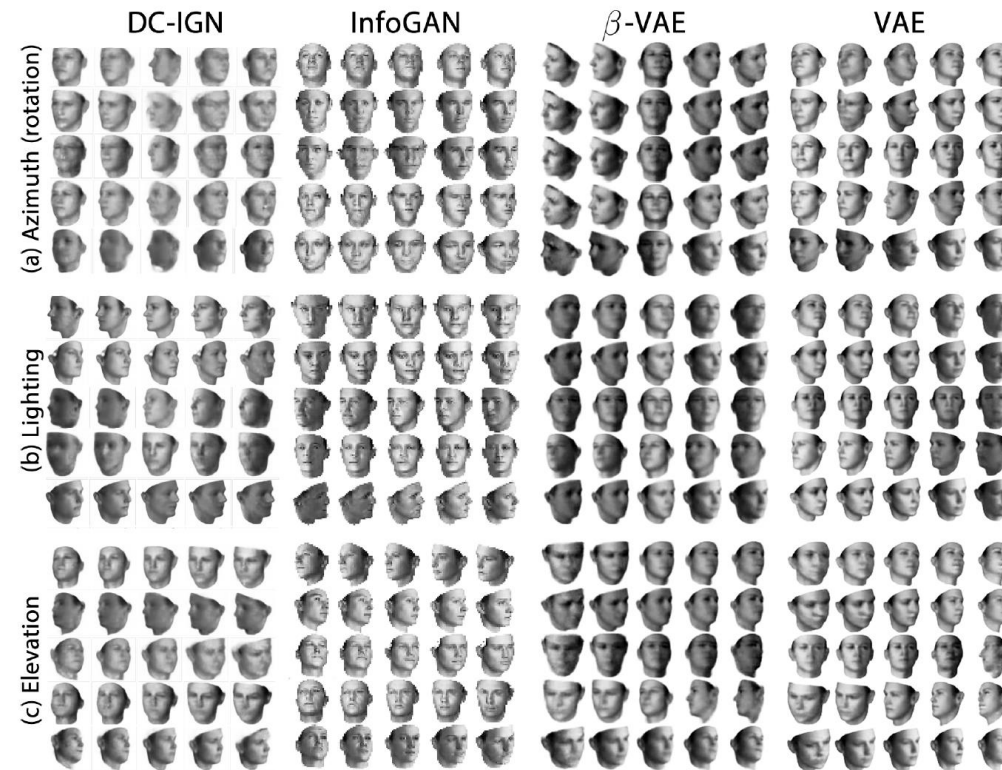


# VAE Variants

- $\beta$ -VAE (Higgins et. al, DeepMind, ICLR 2017)

$$J(\theta, \phi; x) = \underbrace{E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)]}_{\text{Reconstruction}} - \underbrace{\beta KL(q(z|x; \phi) || p(z; \theta))}_{\text{KL penalty}}$$

- Learned factors in  $z$




# VAE Variants

- $\beta$ -VAE (Higgins et. al, DeepMind, ICLR 2017)

$$J(\theta, \phi; x) = \underbrace{E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)]}_{\text{Reconstruction}} - \underbrace{\beta KL(q(z|x; \phi) || p(z; \theta))}_{\text{KL penalty}}$$

- Learned factors in  $z$ 
  - Trade-off between reconstruction and disentangle features!



$\beta$  can be critical!

# VAE Variants

- Understanding disentangling in  $\beta$ -VAE (DeepMind, NIPS 2017)

$$J(\theta, \phi; x) = \underbrace{E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)]}_{\text{Reconstruction}} - \underbrace{\beta}_{\text{KL penalty}} | \underbrace{KL(q(z|x; \phi) || p(z; \theta))}_{\text{KL penalty}} - \underbrace{C}_{\text{Controlled capacity}} |$$


Reconstruction

KL penalty

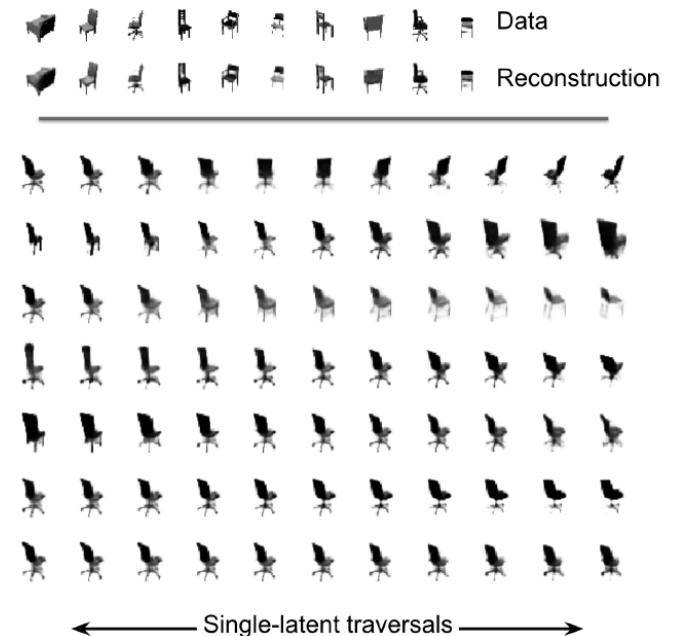
Controlled capacity

- Learned factors in  $z$ 
  - Gradually *increase*  $C$ !

(a) Coloured dSprites



(b) 3D Chairs



# VAE Variants

- $\beta$ -VAE (Higgins et. al, DeepMind, ICLR 2017)

$$J(\theta, \phi; x) = \underbrace{E_{z \sim q(z|x; \phi)} [\log p(x|z; \theta)]}_{\text{Reconstruction}} - \underbrace{\beta KL(q(z|x; \phi) || p(z; \theta))}_{\text{KL penalty}}$$

- Learned factors in  $z$ 
  - A popular (unsupervised) approach for pretraining features
- No free lunch!
  - Challenging Common Assumptions in the Unsupervised Learning of Disentangled Representations, (Google Brain, ICML2019)
  - Disentangle features are fundamentally **impossible** without supervision or model inductive bias
    - Inductive bias or supervision is important (structured model)
    - Empirical successes can be highly random ...
    - **Tune your model hard!**

# Summary

- Generative Model
  - Learn a probability distribution  $p(x; \theta)$
  - Energy-based model:  $p(x) = \frac{1}{Z} \exp(-E(x; \theta))$
  - Flow model:  $x = f(z; \theta)$  ( $f$  is a bijection)
  - Latent variable model:  $p(x, z) = p(x|z)p(z)$
- Variational Autoencoder
  - A computation-efficient design of  $p(x, z)$ 
    - Isomorphic Gaussian wherever possible
  - Variational inference for efficient and stable learning
    - ELBO & reparameterization trick
  - Flexible framework with nice mathematical property
    - But may suffer from blurry outputs... (next lecture!)

Lunch Time